

# Chapter 3

## Data Analysis with Python

**Note :** Will attach all the PDF version of the notebooks at the end of this chapter

### Python Packages for Data Science

- A Python library is a collection of functions and methods that allow you to perform lots of actions without writing any code. The libraries usually contain built in modules providing different functionalities which you can use directly.
- Broad Classifications
  - Scientific Computing Libraries
    - \* **Pandas** offers data structure and tools for effective data manipulation and analysis. It provides facts, access to structured data.
    - \* **NumPy** library uses arrays for its inputs and outputs. It can be extended to objects for matrices and with minor coding changes, developers can perform fast array processing.
    - \* **SciPy** includes functions for some advanced math problems as listed on this slide, as well as data visualization.
  - Data Visualization
    - \* Matplotlib package is the most well known library for data visualization. It is great for making graphs and plots. The graphs are also highly customizable.
    - \* Seaborn which is based on Matplotlib. It's very easy to generate various plots such as heat maps, time series and violin plots.

- Machine Learning
  - \* Scikit-learn library contains tools statistical modeling, including regression, classification, clustering, and so on. This library is built on NumPy, SciPy and Matplotlib.
  - \* Statsmodels is also a Python module that allows users to explore data, estimate statistical models, and perform statistical tests.

## Importing and Exporting Data in Python

---

```
# Importing the required modules
import pandas as pd
url = "something.csv"
df = pd.read_csv(url, header = None) # read csv ( Also xlsx,
    html and so on can be included change read_(type))
df.head(n) # display first n rows
df.tail(n) # display last n rows
path = "something1.csv"
df.to_csv(path) # to save to a csv file
headers = [some list]
df.columns = headers # replaces default headers
df.info # concise summary of dataframe
```

---

## Getting started analysing data in python

---

```
# Importing the required modules
import pandas as pd
df = read_csv(url)
df.dtypes # gives the data types
df.describe() # returns a statistical summary
df.describe(include = "all") # even for object type arguments
```

---

## Accessing databases with Python

- Databases are powerful tools for data scientists
- Python program communicates with the DBMS using API.

- The application program begins its database access with one or more API calls that connect the program to the DBMS.
- To send the SQL statement to the DBMS, the program builds the statement as a text string in a buffer and then makes an API call to pass the buffer contents to the DBMS.
- The application program makes API calls to check the status of its DBMS request and to handle errors. The application program ends its database access with an API call that disconnects it from the database.
- DB-API is Python's standard API for accessing relational databases. It is a standard that allows you to write a single program that works with multiple kinds of relational databases instead of writing a separate program for each one.
- So, if you learn the DB-API functions, then you can apply that knowledge to use any database with Python. The two main concepts in the Python DB-API are connection objects and query objects.
- You use connection objects to connect to a database and manage your transactions. Cursor objects are used to run queries. You open a cursor object and then run queries. The cursor works similar to a cursor in a text processing system where you scroll down in your result set and get your data into the application.
- Cursors are used to scan through the results of a database. Here are the methods used with connection objects. The `cursor()` method returns a new cursor object using the connection.
- The `commit()` method is used to commit any pending transaction to the database.
- The `rollback()` method causes the database to roll back to the start of any pending transaction.
- The `close()` method is used to close a database connection. Let's walk through a Python application that uses the DB-API to query a database.
- First, you import your database module by using the connect API from that module. To open a connection to the database, you use the connection function and pass in the parameters that is the database name, username, and password.

- The connect function returns connection object. After this, you create a cursor object on the connection object. The cursor is used to run queries and fetch results.
- After running the queries using the cursor, we also use the cursor to fetch the results of the query. Finally, when the system is done running the queries, it frees all resources by closing the connection.
- Remember that it is always important to close connections to avoid unused connections taking up resources.

---

```
from dbmodule import connect

#create connection object
connection = connect('databasename', 'username', 'pswd')

#create a cursor object
cursor = connection.cursor()

#Run queries
cursor.execute('select * from mytable')
results = cursor.fetchall()

#free resources
cursor.close()
connection.close()
```

---

## Data Pre-Processing

It is the process of converting or mapping data from one raw form into another format to make it ready for further analysis. Data preprocessing is often called data cleaning or data wrangling, and there are likely other terms

---

```
# Importing the required modules
import pandas as pd
df['column'] # print the column
df['column'] = df['column'] + 1 # adds 1 to each column (Similarly
    a lot of simple arithmetic operations can be done)
```

---

## Dealing with Missing Values in Python

- The first is to check if the person or group that collected the data can go back and find what the actual value should be. Another possibility is just to remove the data where that missing value is found. When you drop data, you could either drop the whole variable or just the single data entry with the missing value. If you don't have a lot of observations with missing data, usually dropping the particular entry is the best. If you're removing data, you want to look to do something that has the least amount of impact.
- Replacing data is better since no data is wasted. However, it is less accurate since we need to replace missing data with a guess of what the data should be.
- One standard for placement technique is to replace missing values by the average value of the entire variable.

---

```
# Importing the required modules
import pandas as pd
import numpy as np

# Imagine if we wanna remove the price column
df.dropna(subset = ["price"],axis = 1, inplace = True) # axis = 1
# drops the entire column

# Imagine if we want to remove a certain missing value in rows of a
# specific(Price column will be considered here) column
df.dropna(subset = ["price"],axis = 0, inplace = True) # axis = 0
# drops the entire row

# To replace missing values
mean = df["price"].mean() # calculate the mean of the price column
df["price"].replace(np.nan,mean) # use the replace function that
# is df["column"].replace(missing_value, new_value) / np.nan =
# find all NAN values

# Maximum value for feature scaling is 1 and minimum is 0
```

---

## Binning

- Binning is when you group values together into bins.

- ```
# Importing the required modules
import pandas as pd
bins = np.linspace(min(df["price"]),max(df["price"]),4) # Returns
    the array bins that contains 4 equally spaced numbers over the
    specified interval of the price.

"""
np.linspace() = returns an array refer to numpy documentation
bins = 3 in this example
But in the argument 4 is given as bins + 1 always contains the
    intervals
Example : 1-3 , 3-4, 4-6
3 bins but 4 numbers
"""

names = ["low","Medium","High"] # group names

df["price-binned"] =pd.cut(df["price"],bins,labels = names,
    include_lowest = True) # df["price-binned"] changes the column
    price to price-binned
```

- Statistical models cannot take in objects or strings as input and for model training only take the numbers as inputs.

```
# Importing the required modules
import pandas as pd
df = read_csv(url)

# Convert categorical variables to dummy variables 0-1
pd.get_dummies(df["fuel"]) # The get_dummies method automatically
    generates a list of numbers, each one corresponding to a
    particular category of the variable.
```

```
df_dummies = pd.get_dummies(df, prefix='Fuel', prefix_sep='.',
columns=['Fuel']) # This is a more advanced method read more in
the docs
```

---

## Categorical Variables into Quantitative variables

- Statistical models cannot take in objects or strings as input and for model training only take the numbers as inputs.

---

```
# Importing the required modules
import pandas as pd
df = read_csv(url)

# Convert catgeorical variables to dummy variables 0-1
pd.get_dummies(df["fuel"]) # The get_dummies method automatically
generates a list of numbers, each one corresponding to a
particular category of the variable.

df_dummies = pd.get_dummies(df, prefix='Fuel', prefix_sep='.',
columns=['Fuel']) # This is a more advanced method read more in
the docs
```

---

## Exploratory Data Analysis

**Note :** This section required visual graphs and plots so the notebook will cover most of it

---

```
# Importing the required modules
import pandas as pd
import seaborn as sns

df.value_counts() # summarize the categorical data

"""BOX PLOTS
1 - Box plots are great way to visualize numeric data, since you
can visualize the various distributions of the data.
2 - Box plots, you can easily spot outliers and also see the
distribution and skewness of the data. Box plots make it easy
to compare between groups.
3 - Main features - Median, IQR, Mean, outliers
```

```

"""
sns.boxplot(x = "drive",y = "price", data = df) # using seaborn
plot a boxplot

""" SCATTER PLOTS
It shows the relationship between two variables.
1 - Typically set the predictor variable on the x-axis or
horizontal axis - independent
2 - Set the target variable on the y-axis or vertical axis -
dependent
"""
# Example

y = df["price"]
x = df["engine-size"]
plt.scatter(x,y)

plt.title("Scatterplot of Engine Size vs Price")
plt.xlabel("Engine Size")
plt.ylabel("Price")

```

---

## Groupby

- In Pandas, this can be done using the group by method. The group by method is used on categorical variables, groups the data into subsets according to the different categories of that variable.
- You can group by a single variable or you can group by multiple variables by passing in multiple variable names.

---

```

# Importing the required modules
import pandas as pd
df = read_csv(url)

df_group = df[["price","size"]]
df_grp = df_group.groupby(["price","size"],as_index =
False).mean() # groupby these columns and only the average
price is returned for each size
df_grp

# Pivot table
# One variable displayed along the columns and the other displayed
along the rows

```



```

df_pivot = df_grp.pivot(index = "drive-wheels", columns
                        ="body-size") # returns a rectangular grid which is easier to
                        understand

# Heatmap
# plot target over multiple variables

plt.pcolor(df_pivot,cmap = 'RdBu') # pass in the data and also
    cmap specifies the color correlation

plt.colorbar()
plt.show()

```

---

## Correlation

- Correlation is a statistical metric for measuring to what extent different variables are interdependent. In other words, when we look at two variables over time, if one variable changes how does this affect change in the other variable?
- correlation doesn't imply causation.
- Pearson Correlation - Pearson correlation method will give you two values: the correlation coefficient and the P-value
- A value close to 1 implies a large positive correlation, while a value close to negative 1 implies a large negative correlation, and a value close to zero implies no correlation between the variables.
- Next, the P-value will tell us how certain we are about the correlation that we calculated. For the P-value, a value less than .001 gives us a strong certainty about the correlation coefficient that we calculated. A value between .001 and .05 gives us moderate certainty. A value between .05 and .1 will give us a weak certainty. And a P-value larger than .1 will give us no certainty of correlation at all. We can say that there is a strong correlation when the correlation coefficient is close to 1 or negative 1, and the P-value is less than .001.

---

```

# Importing the required modules
import pandas as pd
import seaborn as sns

```

```

df = read_csv(url)

# Correlation
sns.regplot(x = "size",y = "price" , data = df)
plt.ylim(0,)

# Pearson Correlation
pearson_coef, p_value =
    stats.pearsonr(df['horsepower'],df['price'])

```

---

## Analysis of Variance ANOVA

- The correlation among different categories.
- The ANOVA test returns two values, the F-test score and the p-value. The F-test calculates the ratio of variation between groups mean, over the variation within each of the sample groups. The p-value shows whether the obtained result is statistically significant.
- F-test calculates the ratio of variation between groups means over the variation within each of the sample group means.

---

```

# Importing the required modules
import pandas as pd
import seaborn as sns
import scipy as stats
df = read_csv(url)

# ANOVA
df_anova = df[["make","price"]]

grouped_anova = df_anova.groupby(["make"]) # return a dataframe
with grouped columns

results =
    stats.f_oneway(grouped_anova.get_group("honda")["price"],grouped_anova.get_group
# Looks Complicated but it is not it just returns the F value and
p- value

```

---

## Linear Regression and Multiple Linear Regression

- Linear regression will refer to one independent variable to make a prediction.
- $y' = m_1x + c$
- Multiple linear regression will refer to multiple independent variables to make a prediction.
- The predictor independent variable  $x$  and the target dependent variable  $y$ .
- Multiple linear regression is used to explain the relationship between one continuous target  $y$  variable and two or more predictor  $x$  variables.
- $y' = m_1x_1 + m_2x_1 + c$

---

```
# Importing the required modules
import pandas as pd
import seaborn as sns
from sklearn.linear_model import LinearRegression

lm = LinearRegression() # create the model
x = df[["highway"]]
y = df[["price"]]

lm.fit(x,y) # fit the model

Yhat = lm.predict(x) # get the prediction

# Multiple Linear Model Estimator

z = df[["horsepower","weight"]]

lm.fit(z,df["price"])

Yhat = lm.predict(x) # returns an array of values

lm.coef_ # returns the slope

lm.intercept_ # returns the intercept
```

---

## Polynomial Regression and Pipelines

- he polynomial regression. We transform our data into a polynomial, then use linear regression to fit the parameter.
- This method is beneficial for describing curvilinear relationships. What is a curvilinear relationship? It's what you get by squaring or setting higher order terms of the predictor variables in the model transforming the data. The model can be quadratic, which means that the predictor variable in the model is squared.
- $y' = b_0 + b_1x + b_2(x^2)$
- It can go upto any order
- 

---

```
# Importing the required modules
import pandas as pd
import seaborn as sns
import scipy as stats
df = read_csv(url)

# Polynomial Regression

f = np.polyfit(x,y,2)
p = np.polyld(f)

print(p)
```

---

## Measures for In-Sample Evaluation

- These measures are a way to numerically determine how good the model fits on our data. Two important measures that we often use to determine the fit of a model are: Mean Square Error (MSE), and R-squared.

## Prediction and Decision Making

- How can we determine if our model is correct? The first thing you should do is make sure your model results make sense. You should always use visualization, numerical measures for evaluation and comparing between different models.

## Model Evaluation and Refinement

**Note :** As this heavily math based and code everything will be explained through the notebook attached below

# data-wrangling

June 21, 2020

Data Wrangling

Welcome!

By the end of this notebook, you will have learned the basics of Data Wrangling!

What is the purpose of Data Wrangling?

Data Wrangling is the process of converting data from the initial format to a format that may be better for analysis.

What is the fuel consumption (L/100k) rate for the diesel car?

Import data

You can find the “Automobile Data Set” from the following link: <https://archive.ics.uci.edu/ml/machine-learning-databases/autos/imports-85.data>. We will be using this data set throughout this course.

Import pandas

```
[57]: import pandas as pd
import matplotlib.pyplot as plt
```

Reading the data set from the URL and adding the related headers.

URL of the dataset

This dataset was hosted on IBM Cloud object click [HERE](#) for free storage

```
[58]: filename = "https://s3-api.us-geo.objectstorage.softlayer.net/cf-courses-data/
↪CognitiveClass/DA0101EN/auto.csv"
```

Python list headers containing name of headers

```
[59]: headers = ["symboling", "normalized-losses", "make", "fuel-type", "aspiration",
↪ "num-of-doors", "body-style",
↪ "drive-wheels", "engine-location", "wheel-base",
↪ "length", "width", "height", "curb-weight", "engine-type",
↪ "num-of-cylinders",
↪ "engine-size", "fuel-system", "bore", "stroke", "compression-ratio", "horsepower",
↪ "peak-rpm", "city-mpg", "highway-mpg", "price"]
```

Use the Pandas method `read_csv()` to load the data from the web address. Set the parameter “names” equal to the Python list “headers”.

```
[60]: df = pd.read_csv(filename, names = headers)
```

Use the method `head()` to display the first five rows of the dataframe.

```
[61]: # To see what the data set looks like, we'll use the head() method.
df.head()
```

```
[61]:      symboling normalized-losses      make fuel-type aspiration num-of-doors \
0          3          ?  alfa-romero    gas      std          two
1          3          ?  alfa-romero    gas      std          two
2          1          ?  alfa-romero    gas      std          two
3          2        164      audi    gas      std          four
4          2        164      audi    gas      std          four

      body-style drive-wheels engine-location  wheel-base  ...  engine-size  \
0  convertible      rwd      front      88.6  ...      130
1  convertible      rwd      front      88.6  ...      130
2   hatchback      rwd      front      94.5  ...      152
3      sedan      fwd      front      99.8  ...      109
4      sedan      4wd      front      99.4  ...      136

      fuel-system  bore  stroke  compression-ratio  horsepower  peak-rpm  city-mpg  \
0      mpfi  3.47    2.68           9.0          111      5000      21
1      mpfi  3.47    2.68           9.0          111      5000      21
2      mpfi  2.68    3.47           9.0          154      5000      19
3      mpfi  3.19    3.40          10.0          102      5500      24
4      mpfi  3.19    3.40           8.0          115      5500      18

      highway-mpg  price
0          27  13495
1          27  16500
2          26  16500
3          30  13950
4          22  17450
```

```
[5 rows x 26 columns]
```

As we can see, several question marks appeared in the dataframe; those are missing values which may hinder our further analysis.

So, how do we identify all those missing values and deal with them?

How to work with missing data?

Steps for working with missing data:

identify missing data

deal with missing data

correct data format

Identify and handle missing values

Identify missing values

Convert “?” to NaN

In the car dataset, missing data comes with the question mark “?”. We replace “?” with NaN (Not a Number), which is Python’s default missing value marker, for reasons of computational speed and convenience. Here we use the function:

to replace A by B

```
[62]: import numpy as np

# replace "?" to NaN
df.replace("?", np.nan, inplace = True)
df.head(5)
```

```
[62]:      symboling normalized-losses      make fuel-type aspiration num-of-doors \
0          3           NaN alfa-romero      gas      std          two
1          3           NaN alfa-romero      gas      std          two
2          1           NaN alfa-romero      gas      std          two
3          2          164      audi      gas      std          four
4          2          164      audi      gas      std          four

      body-style drive-wheels engine-location  wheel-base  ...  engine-size  \
0  convertible      rwd      front      88.6  ...      130
1  convertible      rwd      front      88.6  ...      130
2   hatchback      rwd      front      94.5  ...      152
3      sedan      fwd      front      99.8  ...      109
4      sedan      4wd      front      99.4  ...      136

      fuel-system  bore  stroke compression-ratio horsepower  peak-rpm city-mpg  \
0      mpfi  3.47   2.68           9.0      111      5000      21
1      mpfi  3.47   2.68           9.0      111      5000      21
2      mpfi  2.68   3.47           9.0      154      5000      19
3      mpfi  3.19   3.40          10.0      102      5500      24
4      mpfi  3.19   3.40           8.0      115      5500      18

      highway-mpg  price
0          27  13495
1          27  16500
2          26  16500
3          30  13950
4          22  17450
```



[5 rows x 26 columns]

identify\_missing\_values

Evaluating for Missing Data

The missing values are converted to Python's default. We use Python's built-in functions to identify these missing values. There are two methods to detect missing data:

`.isnull()`

`.notnull()`

The output is a boolean value indicating whether the value that is passed into the argument is in fact missing data.

```
[63]: missing_data = df.isnull()
      missing_data.head(5)
```

```
[63]:  symboling  normalized-losses  make  fuel-type  aspiration  num-of-doors  \
0      False                True  False      False      False      False
1      False                True  False      False      False      False
2      False                True  False      False      False      False
3      False                False  False      False      False      False
4      False                False  False      False      False      False

      body-style  drive-wheels  engine-location  wheel-base  ...  engine-size  \
0      False      False      False      False  ...      False
1      False      False      False      False  ...      False
2      False      False      False      False  ...      False
3      False      False      False      False  ...      False
4      False      False      False      False  ...      False

      fuel-system  bore  stroke  compression-ratio  horsepower  peak-rpm  \
0      False  False  False      False      False      False
1      False  False  False      False      False      False
2      False  False  False      False      False      False
3      False  False  False      False      False      False
4      False  False  False      False      False      False

      city-mpg  highway-mpg  price
0      False      False  False
1      False      False  False
2      False      False  False
3      False      False  False
4      False      False  False
```

[5 rows x 26 columns]

“True” stands for missing value, while “False” stands for not missing value.

Count missing values in each column

Using a for loop in Python, we can quickly figure out the number of missing values in each column. As mentioned above, “True” represents a missing value, “False” means the value is present in the dataset. In the body of the for loop the method “.value\_counts()” counts the number of “True” values.

```
[64]: for column in missing_data.columns.values.tolist():  
      print(column)  
      print (missing_data[column].value_counts())  
      print("")
```

symboling

False 205

Name: symboling, dtype: int64

normalized-losses

False 164

True 41

Name: normalized-losses, dtype: int64

make

False 205

Name: make, dtype: int64

fuel-type

False 205

Name: fuel-type, dtype: int64

aspiration

False 205

Name: aspiration, dtype: int64

num-of-doors

False 203

True 2

Name: num-of-doors, dtype: int64

body-style

False 205

Name: body-style, dtype: int64

drive-wheels

False 205

Name: drive-wheels, dtype: int64

engine-location

False 205

Name: engine-location, dtype: int64

wheel-base

False 205

Name: wheel-base, dtype: int64

length

False 205

Name: length, dtype: int64

width

False 205

Name: width, dtype: int64

height

False 205

Name: height, dtype: int64

curb-weight

False 205

Name: curb-weight, dtype: int64

engine-type

False 205

Name: engine-type, dtype: int64

num-of-cylinders

False 205

Name: num-of-cylinders, dtype: int64

engine-size

False 205

Name: engine-size, dtype: int64

fuel-system

False 205

Name: fuel-system, dtype: int64

bore

False 201

True 4

Name: bore, dtype: int64

stroke

False 201

True 4

Name: stroke, dtype: int64

```
compression-ratio
False      205
Name: compression-ratio, dtype: int64
```

```
horsepower
False      203
True        2
Name: horsepower, dtype: int64
```

```
peak-rpm
False      203
True        2
Name: peak-rpm, dtype: int64
```

```
city-mpg
False      205
Name: city-mpg, dtype: int64
```

```
highway-mpg
False      205
Name: highway-mpg, dtype: int64
```

```
price
False      201
True        4
Name: price, dtype: int64
```

Based on the summary above, each column has 205 rows of data, seven columns containing missing data:

“normalized-losses”: 41 missing data

“num-of-doors”: 2 missing data

“bore”: 4 missing data

“stroke” : 4 missing data

“horsepower”: 2 missing data

“peak-rpm”: 2 missing data

“price”: 4 missing data

Deal with missing data

How to deal with missing data?

drop data a. drop the whole row b. drop the whole column

replace data a. replace it by mean b. replace it by frequency c. replace it based on other functions

Whole columns should be dropped only if most entries in the column are empty. In our dataset, none of the columns are empty enough to drop entirely. We have some freedom in choosing which method to replace data; however, some methods may seem more reasonable than others. We will apply each method to many different columns:

Replace by mean:

“normalized-losses”: 41 missing data, replace them with mean

“stroke”: 4 missing data, replace them with mean

“bore”: 4 missing data, replace them with mean

“horsepower”: 2 missing data, replace them with mean

“peak-rpm”: 2 missing data, replace them with mean

Replace by frequency:

“num-of-doors”: 2 missing data, replace them with “four”.

Reason: 84% sedans is four doors. Since four doors is most frequent, it is most likely to occur

</li>

Drop the whole row:

“price”: 4 missing data, simply delete the whole row

Reason: price is what we want to predict. Any data entry without price data cannot be used for prediction; therefore any row now without price data is not useful to us

</li>

Calculate the average of the column

```
[65]: avg_norm_loss = df["normalized-losses"].astype("float").mean(axis=0)
      print("Average of normalized-losses:", avg_norm_loss)
```

Average of normalized-losses: 122.0

Replace “NaN” by mean value in “normalized-losses” column

```
[66]: df["normalized-losses"].replace(np.nan, avg_norm_loss, inplace=True)
```

Calculate the mean value for ‘bore’ column

```
[67]: avg_bore=df['bore'].astype('float').mean(axis=0)
      print("Average of bore:", avg_bore)
```

Average of bore: 3.3297512437810943

Replace NaN by mean value

```
[68]: df["bore"].replace(np.nan, avg_bore, inplace=True)
```

Question #1:

According to the example above, replace NaN in “stroke” column by mean.

```
[69]: avg_stroke = df["stroke"].astype("float").mean(axis = 0)
      print("Average of stroke:", avg_stroke)
      df["stroke"].replace(np.nan, avg_stroke, inplace = True)
```

Average of stroke: 3.255422885572139

Double-click here for the solution.

Calculate the mean value for the ‘horsepower’ column:

```
[70]: avg_horsepower = df['horsepower'].astype('float').mean(axis=0)
      print("Average horsepower:", avg_horsepower)
```

Average horsepower: 104.25615763546799

Replace “NaN” by mean value:

```
[71]: df['horsepower'].replace(np.nan, avg_horsepower, inplace=True)
```

Calculate the mean value for ‘peak-rpm’ column:

```
[72]: avg_peakrpm=df['peak-rpm'].astype('float').mean(axis=0)
      print("Average peak rpm:", avg_peakrpm)
```

Average peak rpm: 5125.369458128079

Replace NaN by mean value:

```
[73]: df['peak-rpm'].replace(np.nan, avg_peakrpm, inplace=True)
```

To see which values are present in a particular column, we can use the “value\_counts()” method:

```
[74]: df['num-of-doors'].value_counts()
```

```
[74]: four      114
      two       89
      Name: num-of-doors, dtype: int64
```

We can see that four doors are the most common type. We can also use the “idxmax()” method to calculate for us the most common type automatically:

```
[75]: df['num-of-doors'].value_counts().idxmax()
```

```
[75]: 'four'
```

The replacement procedure is very similar to what we have seen previously

```
[76]: #replace the missing 'num-of-doors' values by the most frequent
      df["num-of-doors"].replace(np.nan, "four", inplace=True)
```

Finally, let's drop all rows that do not have price data:

```
[77]: # simply drop whole row with NaN in "price" column
df.dropna(subset=["price"], axis=0, inplace=True)

# reset index, because we dropped two rows
df.reset_index(drop=True, inplace=True)
```

```
[78]: df.head()
```

```
[78]:
```

|   | symboling | normalized-losses | make        | fuel-type | aspiration | num-of-doors | \ |
|---|-----------|-------------------|-------------|-----------|------------|--------------|---|
| 0 | 3         | 122               | alfa-romero | gas       | std        | two          |   |
| 1 | 3         | 122               | alfa-romero | gas       | std        | two          |   |
| 2 | 1         | 122               | alfa-romero | gas       | std        | two          |   |
| 3 | 2         | 164               | audi        | gas       | std        | four         |   |
| 4 | 2         | 164               | audi        | gas       | std        | four         |   |

|   | body-style  | drive-wheels | engine-location | wheel-base | ... | engine-size | \ |
|---|-------------|--------------|-----------------|------------|-----|-------------|---|
| 0 | convertible | rwd          | front           | 88.6       | ... | 130         |   |
| 1 | convertible | rwd          | front           | 88.6       | ... | 130         |   |
| 2 | hatchback   | rwd          | front           | 94.5       | ... | 152         |   |
| 3 | sedan       | fwd          | front           | 99.8       | ... | 109         |   |
| 4 | sedan       | 4wd          | front           | 99.4       | ... | 136         |   |

|   | fuel-system | bore | stroke | compression-ratio | horsepower | peak-rpm | city-mpg | \ |
|---|-------------|------|--------|-------------------|------------|----------|----------|---|
| 0 | mpfi        | 3.47 | 2.68   | 9.0               | 111        | 5000     | 21       |   |
| 1 | mpfi        | 3.47 | 2.68   | 9.0               | 111        | 5000     | 21       |   |
| 2 | mpfi        | 2.68 | 3.47   | 9.0               | 154        | 5000     | 19       |   |
| 3 | mpfi        | 3.19 | 3.40   | 10.0              | 102        | 5500     | 24       |   |
| 4 | mpfi        | 3.19 | 3.40   | 8.0               | 115        | 5500     | 18       |   |

|   | highway-mpg | price |
|---|-------------|-------|
| 0 | 27          | 13495 |
| 1 | 27          | 16500 |
| 2 | 26          | 16500 |
| 3 | 30          | 13950 |
| 4 | 22          | 17450 |

```
[5 rows x 26 columns]
```

Good! Now, we obtain the dataset with no missing values.

Correct data format

We are almost there!

The last step in data cleaning is checking and making sure that all data is in the correct format (int, float, text or other).

In Pandas, we use

.dtype() to check the data type

.astype() to change the data type

Lets list the data types for each column

```
[79]: df.dtypes
```

```
[79]: symboling          int64
normalized-losses      object
make                   object
fuel-type              object
aspiration             object
num-of-doors           object
body-style             object
drive-wheels           object
engine-location        object
wheel-base            float64
length                float64
width                  float64
height                 float64
curb-weight            int64
engine-type            object
num-of-cylinders       object
engine-size            int64
fuel-system            object
bore                   object
stroke                 object
compression-ratio      float64
horsepower             object
peak-rpm              object
city-mpg               int64
highway-mpg            int64
price                  object
dtype: object
```

As we can see above, some columns are not of the correct data type. Numerical variables should have type 'float' or 'int', and variables with strings such as categories should have type 'object'. For example, 'bore' and 'stroke' variables are numerical values that describe the engines, so we should expect them to be of the type 'float' or 'int'; however, they are shown as type 'object'. We have to convert data types into a proper format for each column using the "astype()" method.

```
[80]: df[["bore", "stroke"]] = df[["bore", "stroke"]].astype("float")
df[["normalized-losses"]] = df[["normalized-losses"]].astype("int")
df[["price"]] = df[["price"]].astype("float")
df[["peak-rpm"]] = df[["peak-rpm"]].astype("float")
```

Let us list the columns after the conversion



```
[81]: df.dtypes
```

```
[81]: symboling          int64
normalized-losses      int64
make                   object
fuel-type              object
aspiration             object
num-of-doors           object
body-style             object
drive-wheels           object
engine-location        object
wheel-base            float64
length                float64
width                 float64
height                float64
curb-weight            int64
engine-type            object
num-of-cylinders       object
engine-size            int64
fuel-system            object
bore                  float64
stroke                float64
compression-ratio      float64
horsepower             object
peak-rpm              float64
city-mpg               int64
highway-mpg            int64
price                 float64
dtype: object
```

Wonderful!

Now, we finally obtain the cleaned dataset with no missing values and all data in its proper format.

## Data Standardization

Data is usually collected from different agencies with different formats. (Data Standardization is also a term for a particular type of data normalization, where we subtract the mean and divide by the standard deviation)

What is Standardization?

Standardization is the process of transforming data into a common format which allows the researcher to make the meaningful comparison.

## Example

Transform mpg to L/100km:

In our dataset, the fuel consumption columns “city-mpg” and “highway-mpg” are represented by mpg (miles per gallon) unit. Assume we are developing an application in a country that accept the

fuel consumption with L/100km standard

We will need to apply data transformation to transform mpg into L/100km?

The formula for unit conversion is

$$\text{L/100km} = 235 / \text{mpg}$$

We can do many mathematical operations directly in Pandas.

```
[82]: df.head()
```

```
[82]:      symboling  normalized-losses      make fuel-type aspiration \
0           3           122  alfa-romero      gas      std
1           3           122  alfa-romero      gas      std
2           1           122  alfa-romero      gas      std
3           2           164      audi      gas      std
4           2           164      audi      gas      std

      num-of-doors  body-style drive-wheels engine-location  wheel-base  ... \
0           two  convertible      rwd      front      88.6  ...
1           two  convertible      rwd      front      88.6  ...
2           two   hatchback      rwd      front      94.5  ...
3          four      sedan      fwd      front      99.8  ...
4          four      sedan      4wd      front      99.4  ...

      engine-size  fuel-system  bore  stroke  compression-ratio  horsepower  \
0           130      mpfi  3.47   2.68           9.0           111
1           130      mpfi  3.47   2.68           9.0           111
2           152      mpfi  2.68   3.47           9.0           154
3           109      mpfi  3.19   3.40          10.0           102
4           136      mpfi  3.19   3.40           8.0           115

      peak-rpm  city-mpg  highway-mpg      price
0     5000.0      21      27  13495.0
1     5000.0      21      27  16500.0
2     5000.0      19      26  16500.0
3     5500.0      24      30  13950.0
4     5500.0      18      22  17450.0
```

[5 rows x 26 columns]

```
[83]: # Convert mpg to L/100km by mathematical operation (235 divided by mpg)
df['city-L/100km'] = 235/df["city-mpg"]

# check your transformed data
df.head()
```

```
[83]:      symboling  normalized-losses      make fuel-type aspiration \
0          3          122  alfa-romero      gas      std
1          3          122  alfa-romero      gas      std
2          1          122  alfa-romero      gas      std
3          2          164      audi      gas      std
4          2          164      audi      gas      std

      num-of-doors  body-style drive-wheels engine-location  wheel-base  ... \
0          two  convertible      rwd      front      88.6  ...
1          two  convertible      rwd      front      88.6  ...
2          two   hatchback      rwd      front      94.5  ...
3          four    sedan      fwd      front      99.8  ...
4          four    sedan      4wd      front      99.4  ...

      fuel-system  bore  stroke  compression-ratio  horsepower  peak-rpm  city-mpg \
0          mpfi  3.47   2.68           9.0          111   5000.0      21
1          mpfi  3.47   2.68           9.0          111   5000.0      21
2          mpfi  2.68   3.47           9.0          154   5000.0      19
3          mpfi  3.19   3.40          10.0          102   5500.0      24
4          mpfi  3.19   3.40           8.0          115   5500.0      18

      highway-mpg  price  city-L/100km
0          27  13495.0    11.190476
1          27  16500.0    11.190476
2          26  16500.0    12.368421
3          30  13950.0     9.791667
4          22  17450.0    13.055556

[5 rows x 27 columns]
```

Question #2:

According to the example above, transform mpg to L/100km in the column of “highway-mpg”, and change the name of column to “highway-L/100km”.

```
[84]: df["highway-mpg"] = 235/df["highway-mpg"]
df.rename(columns={"highway-mpg": 'highway-L/100km'}, inplace=True)
df.head()
```

```
[84]:      symboling  normalized-losses      make fuel-type aspiration \
0          3          122  alfa-romero      gas      std
1          3          122  alfa-romero      gas      std
2          1          122  alfa-romero      gas      std
3          2          164      audi      gas      std
4          2          164      audi      gas      std

      num-of-doors  body-style drive-wheels engine-location  wheel-base  ... \
0          two  convertible      rwd      front      88.6  ...
```

|   |      |             |     |       |      |     |
|---|------|-------------|-----|-------|------|-----|
| 1 | two  | convertible | rwd | front | 88.6 | ... |
| 2 | two  | hatchback   | rwd | front | 94.5 | ... |
| 3 | four | sedan       | fwd | front | 99.8 | ... |
| 4 | four | sedan       | 4wd | front | 99.4 | ... |

|   | fuel-system | bore | stroke | compression-ratio | horsepower | peak-rpm | city-mpg | \ |
|---|-------------|------|--------|-------------------|------------|----------|----------|---|
| 0 | mpfi        | 3.47 | 2.68   | 9.0               | 111        | 5000.0   | 21       |   |
| 1 | mpfi        | 3.47 | 2.68   | 9.0               | 111        | 5000.0   | 21       |   |
| 2 | mpfi        | 2.68 | 3.47   | 9.0               | 154        | 5000.0   | 19       |   |
| 3 | mpfi        | 3.19 | 3.40   | 10.0              | 102        | 5500.0   | 24       |   |
| 4 | mpfi        | 3.19 | 3.40   | 8.0               | 115        | 5500.0   | 18       |   |

|   | highway-mpg | price   | city-L/100km |
|---|-------------|---------|--------------|
| 0 | 8.703704    | 13495.0 | 11.190476    |
| 1 | 8.703704    | 16500.0 | 11.190476    |
| 2 | 9.038462    | 16500.0 | 12.368421    |
| 3 | 7.833333    | 13950.0 | 9.791667     |
| 4 | 10.681818   | 17450.0 | 13.055556    |

[5 rows x 27 columns]

Double-click here for the solution.

Data Normalization

Why normalization?

Normalization is the process of transforming values of several variables into a similar range. Typical normalizations include scaling the variable so the variable average is 0, scaling the variable so the variance is 1, or scaling variable so the variable values range from 0 to 1

Example

To demonstrate normalization, let's say we want to scale the columns "length", "width" and "height"

Target: would like to Normalize those variables so their value ranges from 0 to 1.

Approach: replace original value by (original value)/(maximum value)

```
[85]: # replace (original value) by (original value)/(maximum value)
df['length'] = df['length']/df['length'].max()
df['width'] = df['width']/df['width'].max()
```

Questiont #3:

According to the example above, normalize the column "height".

```
[86]: df['height'] = df['height']/df['height'].max()
df[["length", "width", "height"]].head()
```

```
[86]:      length      width      height
0  0.811148  0.890278  0.816054
1  0.811148  0.890278  0.816054
2  0.822681  0.909722  0.876254
3  0.848630  0.919444  0.908027
4  0.848630  0.922222  0.908027
```

Double-click here for the solution.

Here we can see, we've normalized "length", "width" and "height" in the range of [0,1].

Binning

Why binning?

Binning is a process of transforming continuous numerical variables into discrete categorical 'bins', for grouped analysis.

Example:

In our dataset, "horsepower" is a real valued variable ranging from 48 to 288, it has 57 unique values. What if we only care about the price difference between cars with high horsepower, medium horsepower, and little horsepower (3 types)? Can we rearrange them into three 'bins' to simplify analysis?

We will use the Pandas method 'cut' to segment the 'horsepower' column into 3 bins

Example of Binning Data In Pandas

Convert data to correct format


```
[87]: df["horsepower"]=df["horsepower"].astype(int, copy=True)
```

Lets plot the histogram of horsepower, to see what the distribution of horsepower looks like.

```
[88]: %matplotlib inline
import matplotlib as plt
from matplotlib import pyplot
plt.pyplot.hist(df["horsepower"])

# set x/y labels and plot title
plt.pyplot.xlabel("horsepower")
plt.pyplot.ylabel("count")
plt.pyplot.title("horsepower bins")
```

```
[88]: Text(0.5, 1.0, 'horsepower bins')
```



We would like 3 bins of equal size bandwidth so we use numpy's `linspace(start_value, end_value, numbers_generated)` function.

Since we want to include the minimum value of horsepower we want to set `start_value=min(df["horsepower"])`.

Since we want to include the maximum value of horsepower we want to set `end_value=max(df["horsepower"])`.

Since we are building 3 bins of equal length, there should be 4 dividers, so `numbers_generated=4`.

We build a bin array, with a minimum value to a maximum value, with bandwidth calculated above. The bins will be values used to determine when one bin ends and another begins.

```
[89]: bins = np.linspace(min(df["horsepower"]), max(df["horsepower"]), 4)
      bins
```

```
[89]: array([ 48.          , 119.33333333, 190.66666667, 262.          ])
```

We set group names:

```
[90]: group_names = ['Low', 'Medium', 'High']
```

We apply the function "cut" to determine what each value of "df['horsepower']" belongs to.

```
[91]: df['horsepower-binned'] = pd.cut(df['horsepower'], bins, labels=group_names,
    ↪include_lowest=True )
df[['horsepower', 'horsepower-binned']].head(20)
```

```
[91]:      horsepower horsepower-binned
0          111             Low
1          111             Low
2          154           Medium
3          102             Low
4          115             Low
5          110             Low
6          110             Low
7          110             Low
8          140           Medium
9          101             Low
10         101             Low
11         121           Medium
12         121           Medium
13         121           Medium
14         182           Medium
15         182           Medium
16         182           Medium
17          48             Low
18          70             Low
19          70             Low
```

Lets see the number of vehicles in each bin.

```
[92]: df["horsepower-binned"].value_counts()
```


```
[92]: Low          153
      Medium       43
      High         5
      Name: horsepower-binned, dtype: int64
```

Lets plot the distribution of each bin.

```
[93]: %matplotlib inline
import matplotlib as plt
from matplotlib import pyplot
pyplot.bar(group_names, df["horsepower-binned"].value_counts())

# set x/y labels and plot title
plt.pyplot.xlabel("horsepower")
plt.pyplot.ylabel("count")
plt.pyplot.title("horsepower bins")
```

```
[93]: Text(0.5, 1.0, 'horsepower bins')
```



Check the dataframe above carefully, you will find the last column provides the bins for “horsepower” with 3 categories (“Low”, “Medium” and “High”).

We successfully narrow the intervals from 57 to 3!

Bins visualization

Normally, a histogram is used to visualize the distribution of bins we created above.

```
[94]: %matplotlib inline
import matplotlib as plt
from matplotlib import pyplot


a = (0,1,2)

# draw histogram of attribute "horsepower" with bins = 3
plt.pyplot.hist(df["horsepower"], bins = 3)

# set x/y labels and plot title
plt.pyplot.xlabel("horsepower")
plt.pyplot.ylabel("count")
plt.pyplot.title("horsepower bins")
```

```
[94]: Text(0.5, 1.0, 'horsepower bins')
```





The plot above shows the binning result for attribute “horsepower”.

Indicator variable (or dummy variable)

What is an indicator variable?

An indicator variable (or dummy variable) is a numerical variable used to label categories. They are called ‘dummies’ because the numbers themselves don’t have inherent meaning.

Why we use indicator variables?

So we can use categorical variables for regression analysis in the later modules.

Example

We see the column “fuel-type” has two unique values, “gas” or “diesel”. Regression doesn’t understand words, only numbers. To use this attribute in regression analysis, we convert “fuel-type” into indicator variables.

We will use the panda’s method ‘get\_dummies’ to assign numerical values to different categories of fuel type.

```
[95]: df.columns
```

```
[95]: Index(['symboling', 'normalized-losses', 'make', 'fuel-type', 'aspiration',  
          'num-of-doors', 'body-style', 'drive-wheels', 'engine-location',  
          'wheel-base', 'length', 'width', 'height', 'curb-weight', 'engine-type',  
          'num-of-cylinders', 'engine-size', 'fuel-system', 'bore', 'stroke',
```

```
'compression-ratio', 'horsepower', 'peak-rpm', 'city-mpg',
'highway-mpg', 'price', 'city-L/100km', 'horsepower-binned'],
dtype='object')
```

get indicator variables and assign it to data frame “dummy\_variable\_1”

```
[96]: dummy_variable_1 = pd.get_dummies(df["fuel-type"])
dummy_variable_1.head()
```

```
[96]:   diesel  gas
0       0    1
1       0    1
2       0    1
3       0    1
4       0    1
```

change column names for clarity

```
[97]: dummy_variable_1.rename(columns={'fuel-type-diesel':'gas', 'fuel-type-gas':
↳ 'diesel'}, inplace=True)
dummy_variable_1.head()
```

```
[97]:   diesel  gas
0       0    1
1       0    1
2       0    1
3       0    1
4       0    1
```

We now have the value 0 to represent “gas” and 1 to represent “diesel” in the column “fuel-type”. We will now insert this column back into our original dataset.

```
[98]: # merge data frame "df" and "dummy_variable_1"
df = pd.concat([df, dummy_variable_1], axis=1)

# drop original column "fuel-type" from "df"
df.drop("fuel-type", axis = 1, inplace=True)
```

```
[99]: df.head()
```

```
[99]:   symboling  normalized-losses      make aspiration num-of-doors \
0         3             122  alfa-romero      std           two
1         3             122  alfa-romero      std           two
2         1             122  alfa-romero      std           two
3         2             164      audi      std           four
4         2             164      audi      std           four

   body-style drive-wheels engine-location  wheel-base  length  ... \
```

|   |             |     |       |      |          |     |
|---|-------------|-----|-------|------|----------|-----|
| 0 | convertible | rwd | front | 88.6 | 0.811148 | ... |
| 1 | convertible | rwd | front | 88.6 | 0.811148 | ... |
| 2 | hatchback   | rwd | front | 94.5 | 0.822681 | ... |
| 3 | sedan       | fwd | front | 99.8 | 0.848630 | ... |
| 4 | sedan       | 4wd | front | 99.4 | 0.848630 | ... |

|   | compression-ratio | horsepower | peak-rpm | city-mpg | highway-mpg | price \ |
|---|-------------------|------------|----------|----------|-------------|---------|
| 0 | 9.0               | 111        | 5000.0   | 21       | 8.703704    | 13495.0 |
| 1 | 9.0               | 111        | 5000.0   | 21       | 8.703704    | 16500.0 |
| 2 | 9.0               | 154        | 5000.0   | 19       | 9.038462    | 16500.0 |
| 3 | 10.0              | 102        | 5500.0   | 24       | 7.833333    | 13950.0 |
| 4 | 8.0               | 115        | 5500.0   | 18       | 10.681818   | 17450.0 |

|   | city-L/100km | horsepower-binned | diesel | gas |
|---|--------------|-------------------|--------|-----|
| 0 | 11.190476    | Low               | 0      | 1   |
| 1 | 11.190476    | Low               | 0      | 1   |
| 2 | 12.368421    | Medium            | 0      | 1   |
| 3 | 9.791667     | Low               | 0      | 1   |
| 4 | 13.055556    | Low               | 0      | 1   |

[5 rows x 29 columns]

The last two columns are now the indicator variable representation of the fuel-type variable. It's all 0s and 1s now.

Question #4:

As above, create indicator variable to the column of “aspiration”: “std” to 0, while “turbo” to 1.

```
[100]: dummy_variable_2 = pd.get_dummies(df['aspiration'])
dummy_variable_2.rename(columns={'std':'aspiration-std', 'turbo':
→ 'aspiration-turbo'}, inplace=True)
dummy_variable_2.head()
```

```
[100]: aspiration-std aspiration-turbo
0          1          0
1          1          0
2          1          0
3          1          0
4          1          0
```

Double-click here for the solution.

Question #5:

Merge the new dataframe to the original dataframe then drop the column ‘aspiration’

```
[101]: df = pd.concat([df, dummy_variable_2], axis=1)
df.drop('aspiration', axis = 1, inplace=True)
```

Double-click here for the solution.

save the new csv

```
[102]: df.to_csv('clean_df.csv')
```

<p><a href="https://cocl.us/corsera\_da0101en\_notebook\_bottom"><img src="https://s3-api.us-geo.

About the Authors:

This notebook was written by Mahdi Noorian PhD, Joseph Santarcangelo, Bahare Talayian, Eric Xiao, Steven Dong, Parizad, Hima Vsudevan and Fiorella Wenver and Yi Yao.

Joseph Santarcangelo is a Data Scientist at IBM, and holds a PhD in Electrical Engineering. His research focused on using Machine Learning, Signal Processing, and Computer Vision to determine how videos impact human cognition. Joseph has been working for IBM since he completed his PhD.

Copyright © 2018 IBM Developer Skills Network. This notebook and its source code are released under the terms of the MIT License.

# exploratory-data-analysis

June 21, 2020

## Exploratory Data Analysis

Welcome!

In this section, we will explore several methods to see if certain characteristics or features can be used to predict car price.

What are the main characteristics which have the most impact on the car price?

### 1. Import Data from Module 2

Setup

Import libraries

```
[55]: import pandas as pd
import numpy as np
```

load data and store in dataframe df:

This dataset was hosted on IBM Cloud object click [HERE](#) for free storage

```
[56]: path='https://s3-api.us-gio.objectstorage.softlayer.net/cf-courses-data/
↳CognitiveClass/DA0101EN/automobileEDA.csv'
df = pd.read_csv(path)
df.head()
```

```
[56]:   symboling  normalized-losses      make aspiration num-of-doors \
0         3             122  alfa-romero      std           two
1         3             122  alfa-romero      std           two
2         1             122  alfa-romero      std           two
3         2             164      audi      std           four
4         2             164      audi      std           four

   body-style drive-wheels engine-location  wheel-base  length  ... \
0  convertible         rwd         front        88.6  0.811148  ...
1  convertible         rwd         front        88.6  0.811148  ...
2   hatchback         rwd         front        94.5  0.822681  ...
3      sedan         fwd         front        99.8  0.848630  ...
4      sedan         4wd         front        99.4  0.848630  ...

compression-ratio  horsepower  peak-rpm city-mpg highway-mpg  price \
```

|   |      |       |        |    |    |         |
|---|------|-------|--------|----|----|---------|
| 0 | 9.0  | 111.0 | 5000.0 | 21 | 27 | 13495.0 |
| 1 | 9.0  | 111.0 | 5000.0 | 21 | 27 | 16500.0 |
| 2 | 9.0  | 154.0 | 5000.0 | 19 | 26 | 16500.0 |
| 3 | 10.0 | 102.0 | 5500.0 | 24 | 30 | 13950.0 |
| 4 | 8.0  | 115.0 | 5500.0 | 18 | 22 | 17450.0 |

|   | city-L/100km | horsepower-binned | diesel | gas |
|---|--------------|-------------------|--------|-----|
| 0 | 11.190476    | Medium            | 0      | 1   |
| 1 | 11.190476    | Medium            | 0      | 1   |
| 2 | 12.368421    | Medium            | 0      | 1   |
| 3 | 9.791667     | Medium            | 0      | 1   |
| 4 | 13.055556    | Medium            | 0      | 1   |

[5 rows x 29 columns]

## 2. Analyzing Individual Feature Patterns using Visualization

To install seaborn we use the pip which is the python package manager.

```
[57]: %%capture
      ! pip install seaborn
```

Import visualization packages “Matplotlib” and “Seaborn”, don’t forget about “%matplotlib inline” to plot in a Jupyter notebook.

```
[58]: import matplotlib.pyplot as plt
      import seaborn as sns
      %matplotlib inline
```

How to choose the right visualization method?

When visualizing individual variables, it is important to first understand what type of variable you are dealing with. This will help us find the right visualization method for that variable.

```
[59]: # list the data types for each column
      print(df.dtypes)
```

|                   |         |
|-------------------|---------|
| symboling         | int64   |
| normalized-losses | int64   |
| make              | object  |
| aspiration        | object  |
| num-of-doors      | object  |
| body-style        | object  |
| drive-wheels      | object  |
| engine-location   | object  |
| wheel-base        | float64 |
| length            | float64 |
| width             | float64 |
| height            | float64 |

```

curb-weight      int64
engine-type      object
num-of-cylinders object
engine-size      int64
fuel-system      object
bore             float64
stroke           float64
compression-ratio float64
horsepower       float64
peak-rpm         float64
city-mpg         int64
highway-mpg      int64
price            float64
city-L/100km     float64
horsepower-binned object
diesel           int64
gas              int64
dtype: object

```

Question #1:

What is the data type of the column “peak-rpm”?

[Double-click here for the solution.](#)

for example, we can calculate the correlation between variables of type “int64” or “float64” using the method “corr”:

```
[60]: df.corr()
```

```

[60]:      symboling  normalized-losses  wheel-base  length  \
symboling      1.000000          0.466264   -0.535987 -0.365404
normalized-losses  0.466264          1.000000   -0.056661  0.019424
wheel-base     -0.535987         -0.056661    1.000000  0.876024
length         -0.365404          0.019424    0.876024  1.000000
width          -0.242423          0.086802    0.814507  0.857170
height         -0.550160         -0.373737    0.590742  0.492063
curb-weight    -0.233118          0.099404    0.782097  0.880665
engine-size    -0.110581          0.112360    0.572027  0.685025
bore           -0.140019         -0.029862    0.493244  0.608971
stroke         -0.008245          0.055563    0.158502  0.124139
compression-ratio -0.182196         -0.114713    0.250313  0.159733
horsepower      0.075819          0.217299    0.371147  0.579821
peak-rpm       0.279740          0.239543   -0.360305 -0.285970
city-mpg       -0.035527         -0.225016   -0.470606 -0.665192
highway-mpg     0.036233         -0.181877   -0.543304 -0.698142
price          -0.082391          0.133999    0.584642  0.690628
city-L/100km    0.066171          0.238567    0.476153  0.657373
diesel         -0.196735         -0.101546    0.307237  0.211187
gas             0.196735          0.101546   -0.307237 -0.211187

```

|                   | width     | height    | curb-weight | engine-size | bore \    |
|-------------------|-----------|-----------|-------------|-------------|-----------|
| symboling         | -0.242423 | -0.550160 | -0.233118   | -0.110581   | -0.140019 |
| normalized-losses | 0.086802  | -0.373737 | 0.099404    | 0.112360    | -0.029862 |
| wheel-base        | 0.814507  | 0.590742  | 0.782097    | 0.572027    | 0.493244  |
| length            | 0.857170  | 0.492063  | 0.880665    | 0.685025    | 0.608971  |
| width             | 1.000000  | 0.306002  | 0.866201    | 0.729436    | 0.544885  |
| height            | 0.306002  | 1.000000  | 0.307581    | 0.074694    | 0.180449  |
| curb-weight       | 0.866201  | 0.307581  | 1.000000    | 0.849072    | 0.644060  |
| engine-size       | 0.729436  | 0.074694  | 0.849072    | 1.000000    | 0.572609  |
| bore              | 0.544885  | 0.180449  | 0.644060    | 0.572609    | 1.000000  |
| stroke            | 0.188829  | -0.062704 | 0.167562    | 0.209523    | -0.055390 |
| compression-ratio | 0.189867  | 0.259737  | 0.156433    | 0.028889    | 0.001263  |
| horsepower        | 0.615077  | -0.087027 | 0.757976    | 0.822676    | 0.566936  |
| peak-rpm          | -0.245800 | -0.309974 | -0.279361   | -0.256733   | -0.267392 |
| city-mpg          | -0.633531 | -0.049800 | -0.749543   | -0.650546   | -0.582027 |
| highway-mpg       | -0.680635 | -0.104812 | -0.794889   | -0.679571   | -0.591309 |
| price             | 0.751265  | 0.135486  | 0.834415    | 0.872335    | 0.543155  |
| city-L/100km      | 0.673363  | 0.003811  | 0.785353    | 0.745059    | 0.554610  |
| diesel            | 0.244356  | 0.281578  | 0.221046    | 0.070779    | 0.054458  |
| gas               | -0.244356 | -0.281578 | -0.221046   | -0.070779   | -0.054458 |

|                   | stroke    | compression-ratio | horsepower | peak-rpm \ |
|-------------------|-----------|-------------------|------------|------------|
| symboling         | -0.008245 | -0.182196         | 0.075819   | 0.279740   |
| normalized-losses | 0.055563  | -0.114713         | 0.217299   | 0.239543   |
| wheel-base        | 0.158502  | 0.250313          | 0.371147   | -0.360305  |
| length            | 0.124139  | 0.159733          | 0.579821   | -0.285970  |
| width             | 0.188829  | 0.189867          | 0.615077   | -0.245800  |
| height            | -0.062704 | 0.259737          | -0.087027  | -0.309974  |
| curb-weight       | 0.167562  | 0.156433          | 0.757976   | -0.279361  |
| engine-size       | 0.209523  | 0.028889          | 0.822676   | -0.256733  |
| bore              | -0.055390 | 0.001263          | 0.566936   | -0.267392  |
| stroke            | 1.000000  | 0.187923          | 0.098462   | -0.065713  |
| compression-ratio | 0.187923  | 1.000000          | -0.214514  | -0.435780  |
| horsepower        | 0.098462  | -0.214514         | 1.000000   | 0.107885   |
| peak-rpm          | -0.065713 | -0.435780         | 0.107885   | 1.000000   |
| city-mpg          | -0.034696 | 0.331425          | -0.822214  | -0.115413  |
| highway-mpg       | -0.035201 | 0.268465          | -0.804575  | -0.058598  |
| price             | 0.082310  | 0.071107          | 0.809575   | -0.101616  |
| city-L/100km      | 0.037300  | -0.299372         | 0.889488   | 0.115830   |
| diesel            | 0.241303  | 0.985231          | -0.169053  | -0.475812  |
| gas               | -0.241303 | -0.985231         | 0.169053   | 0.475812   |

|                   | city-mpg  | highway-mpg | price     | city-L/100km | diesel \  |
|-------------------|-----------|-------------|-----------|--------------|-----------|
| symboling         | -0.035527 | 0.036233    | -0.082391 | 0.066171     | -0.196735 |
| normalized-losses | -0.225016 | -0.181877   | 0.133999  | 0.238567     | -0.101546 |
| wheel-base        | -0.470606 | -0.543304   | 0.584642  | 0.476153     | 0.307237  |



|                   |           |           |           |           |           |
|-------------------|-----------|-----------|-----------|-----------|-----------|
| length            | -0.665192 | -0.698142 | 0.690628  | 0.657373  | 0.211187  |
| width             | -0.633531 | -0.680635 | 0.751265  | 0.673363  | 0.244356  |
| height            | -0.049800 | -0.104812 | 0.135486  | 0.003811  | 0.281578  |
| curb-weight       | -0.749543 | -0.794889 | 0.834415  | 0.785353  | 0.221046  |
| engine-size       | -0.650546 | -0.679571 | 0.872335  | 0.745059  | 0.070779  |
| bore              | -0.582027 | -0.591309 | 0.543155  | 0.554610  | 0.054458  |
| stroke            | -0.034696 | -0.035201 | 0.082310  | 0.037300  | 0.241303  |
| compression-ratio | 0.331425  | 0.268465  | 0.071107  | -0.299372 | 0.985231  |
| horsepower        | -0.822214 | -0.804575 | 0.809575  | 0.889488  | -0.169053 |
| peak-rpm          | -0.115413 | -0.058598 | -0.101616 | 0.115830  | -0.475812 |
| city-mpg          | 1.000000  | 0.972044  | -0.686571 | -0.949713 | 0.265676  |
| highway-mpg       | 0.972044  | 1.000000  | -0.704692 | -0.930028 | 0.198690  |
| price             | -0.686571 | -0.704692 | 1.000000  | 0.789898  | 0.110326  |
| city-L/100km      | -0.949713 | -0.930028 | 0.789898  | 1.000000  | -0.241282 |
| diesel            | 0.265676  | 0.198690  | 0.110326  | -0.241282 | 1.000000  |
| gas               | -0.265676 | -0.198690 | -0.110326 | 0.241282  | -1.000000 |

|                   | gas       |
|-------------------|-----------|
| symboling         | 0.196735  |
| normalized-losses | 0.101546  |
| wheel-base        | -0.307237 |
| length            | -0.211187 |
| width             | -0.244356 |
| height            | -0.281578 |
| curb-weight       | -0.221046 |
| engine-size       | -0.070779 |
| bore              | -0.054458 |
| stroke            | -0.241303 |
| compression-ratio | -0.985231 |
| horsepower        | 0.169053  |
| peak-rpm          | 0.475812  |
| city-mpg          | -0.265676 |
| highway-mpg       | -0.198690 |
| price             | -0.110326 |
| city-L/100km      | 0.241282  |
| diesel            | -1.000000 |
| gas               | 1.000000  |

The diagonal elements are always one; we will study correlation more precisely Pearson correlation in-depth at the end of the notebook.

Question #2:

Find the correlation between the following columns: bore, stroke, compression-ratio , and horsepower.

Hint: if you would like to select those columns use the following syntax: `df[['bore','stroke','compression-ratio','horsepower']]`

```
[61]: df[['bore', 'stroke', 'compression-ratio', 'horsepower']].corr()
```

```
[61]:
```

|                   | bore      | stroke    | compression-ratio | horsepower |
|-------------------|-----------|-----------|-------------------|------------|
| bore              | 1.000000  | -0.055390 | 0.001263          | 0.566936   |
| stroke            | -0.055390 | 1.000000  | 0.187923          | 0.098462   |
| compression-ratio | 0.001263  | 0.187923  | 1.000000          | -0.214514  |
| horsepower        | 0.566936  | 0.098462  | -0.214514         | 1.000000   |

Double-click here for the solution.

Continuous numerical variables:

Continuous numerical variables are variables that may contain any value within some range. Continuous numerical variables can have the type “int64” or “float64”. A great way to visualize these variables is by using scatterplots with fitted lines.

In order to start understanding the (linear) relationship between an individual variable and the price. We can do this by using “regplot”, which plots the scatterplot plus the fitted regression line for the data.


Let’s see several examples of different linear relationships:

Positive linear relationship

Let’s find the scatterplot of “engine-size” and “price”

```
[62]: # Engine size as potential predictor variable of price
sns.regplot(x="engine-size", y="price", data=df)
plt.ylim(0,)
```

```
[62]: (0, 56053.02835800785)
```



As the engine-size goes up, the price goes up: this indicates a positive direct correlation between these two variables. Engine size seems like a pretty good predictor of price since the regression line is almost a perfect diagonal line.

We can examine the correlation between 'engine-size' and 'price' and see it's approximately 0.87


```
[63]: df[["engine-size", "price"]].corr()
```

```
[63]:      engine-size    price
engine-size    1.000000  0.872335
price          0.872335  1.000000
```

Highway mpg is a potential predictor variable of price

```
[64]: sns.regplot(x="highway-mpg", y="price", data=df)
```

```
[64]: <matplotlib.axes._subplots.AxesSubplot at 0x7f4fb5a5b908>
```



As the highway-mpg goes up, the price goes down: this indicates an inverse/negative relationship between these two variables. Highway mpg could potentially be a predictor of price.

We can examine the correlation between 'highway-mpg' and 'price' and see it's approximately -0.704

```
[65]: df[['highway-mpg', 'price']].corr()
```


```
[65]:      highway-mpg    price
highway-mpg    1.000000 -0.704692
price         -0.704692  1.000000
```

Weak Linear Relationship

Let's see if "Peak-rpm" as a predictor variable of "price".

```
[66]: sns.regplot(x="peak-rpm", y="price", data=df)
```

```
[66]: <matplotlib.axes._subplots.AxesSubplot at 0x7f4fb59cc2e8>
```



Peak rpm does not seem like a good predictor of the price at all since the regression line is close to horizontal. Also, the data points are very scattered and far from the fitted line, showing lots of variability. Therefore it's not a reliable variable.

We can examine the correlation between 'peak-rpm' and 'price' and see it's approximately -0.101616

```
[67]: df[['peak-rpm', 'price']].corr()
```

```
[67]:      peak-rpm    price
peak-rpm  1.000000 -0.101616
price    -0.101616  1.000000
```

Question 3 a):

Find the correlation between x="stroke", y="price".

Hint: if you would like to select those columns use the following syntax: `df[["stroke", "price"]]`

```
[68]: df[["stroke", "price"]].corr()
```

```
[68]:      stroke    price
stroke  1.00000  0.08231
price   0.08231  1.00000
```


[Double-click here for the solution.](#)

Question 3 b):

Given the correlation results between “price” and “stroke” do you expect a linear relationship? Verify your results using the function “regplot()”.

```
[69]: sns.regplot(x="stroke", y="price", data=df)
```

```
[69]: <matplotlib.axes._subplots.AxesSubplot at 0x7f4fb593cda0>
```



[Double-click here for the solution.](#)


Categorical variables

These are variables that describe a ‘characteristic’ of a data unit, and are selected from a small group of categories. The categorical variables can have the type “object” or “int64”. A good way to visualize categorical variables is by using boxplots.

Let’s look at the relationship between “body-style” and “price”.

```
[70]: sns.boxplot(x="body-style", y="price", data=df)
```


```
[70]: <matplotlib.axes._subplots.AxesSubplot at 0x7f4fb58b46d8>
```



We see that the distributions of price between the different body-style categories have a significant overlap, and so body-style would not be a good predictor of price. Let's examine engine "engine-location" and "price":

```
[71]: sns.boxplot(x="engine-location", y="price", data=df)
```

```
[71]: <matplotlib.axes._subplots.AxesSubplot at 0x7f4fb5919358>
```




Here we see that the distribution of price between these two engine-location categories, front and rear, are distinct enough to take engine-location as a potential good predictor of price.

Let's examine "drive-wheels" and "price".

```
[72]: # drive-wheels
sns.boxplot(x="drive-wheels", y="price", data=df)
```

```
[72]: <matplotlib.axes._subplots.AxesSubplot at 0x7f4f842aeba8>
```





Here we see that the distribution of price between the different drive-wheels categories differs; as such drive-wheels could potentially be a predictor of price.

### 3. Descriptive Statistical Analysis

Let's first take a look at the variables by utilizing a description method.

The describe function automatically computes basic statistics for all continuous variables. Any NaN values are automatically skipped in these statistics.

This will show:

the count of that variable

the mean

the standard deviation (std)

the minimum value

the IQR (Interquartile Range: 25%, 50% and 75%)

the maximum value

We can apply the method “describe” as follows:

```
[73]: df.describe()
```

```
[73]:      symboling  normalized-losses  wheel-base    length    width  \
count    201.000000         201.00000  201.000000  201.000000  201.000000
```

|      |           |           |            |          |          |
|------|-----------|-----------|------------|----------|----------|
| mean | 0.840796  | 122.00000 | 98.797015  | 0.837102 | 0.915126 |
| std  | 1.254802  | 31.99625  | 6.066366   | 0.059213 | 0.029187 |
| min  | -2.000000 | 65.00000  | 86.600000  | 0.678039 | 0.837500 |
| 25%  | 0.000000  | 101.00000 | 94.500000  | 0.801538 | 0.890278 |
| 50%  | 1.000000  | 122.00000 | 97.000000  | 0.832292 | 0.909722 |
| 75%  | 2.000000  | 137.00000 | 102.400000 | 0.881788 | 0.925000 |
| max  | 3.000000  | 256.00000 | 120.900000 | 1.000000 | 1.000000 |

|       |            |             |             |            |            |
|-------|------------|-------------|-------------|------------|------------|
|       | height     | curb-weight | engine-size | bore       | stroke \   |
| count | 201.000000 | 201.000000  | 201.000000  | 201.000000 | 197.000000 |
| mean  | 53.766667  | 2555.666667 | 126.875622  | 3.330692   | 3.256904   |
| std   | 2.447822   | 517.296727  | 41.546834   | 0.268072   | 0.319256   |
| min   | 47.800000  | 1488.000000 | 61.000000   | 2.540000   | 2.070000   |
| 25%   | 52.000000  | 2169.000000 | 98.000000   | 3.150000   | 3.110000   |
| 50%   | 54.100000  | 2414.000000 | 120.000000  | 3.310000   | 3.290000   |
| 75%   | 55.500000  | 2926.000000 | 141.000000  | 3.580000   | 3.410000   |
| max   | 59.800000  | 4066.000000 | 326.000000  | 3.940000   | 4.170000   |

|       |                   |            |             |            |               |
|-------|-------------------|------------|-------------|------------|---------------|
|       | compression-ratio | horsepower | peak-rpm    | city-mpg   | highway-mpg \ |
| count | 201.000000        | 201.000000 | 201.000000  | 201.000000 | 201.000000    |
| mean  | 10.164279         | 103.405534 | 5117.665368 | 25.179104  | 30.686567     |
| std   | 4.004965          | 37.365700  | 478.113805  | 6.423220   | 6.815150      |
| min   | 7.000000          | 48.000000  | 4150.000000 | 13.000000  | 16.000000     |
| 25%   | 8.600000          | 70.000000  | 4800.000000 | 19.000000  | 25.000000     |
| 50%   | 9.000000          | 95.000000  | 5125.369458 | 24.000000  | 30.000000     |
| 75%   | 9.400000          | 116.000000 | 5500.000000 | 30.000000  | 34.000000     |
| max   | 23.000000         | 262.000000 | 6600.000000 | 49.000000  | 54.000000     |

|       |              |              |            |            |
|-------|--------------|--------------|------------|------------|
|       | price        | city-L/100km | diesel     | gas        |
| count | 201.000000   | 201.000000   | 201.000000 | 201.000000 |
| mean  | 13207.129353 | 9.944145     | 0.099502   | 0.900498   |
| std   | 7947.066342  | 2.534599     | 0.300083   | 0.300083   |
| min   | 5118.000000  | 4.795918     | 0.000000   | 0.000000   |
| 25%   | 7775.000000  | 7.833333     | 0.000000   | 1.000000   |
| 50%   | 10295.000000 | 9.791667     | 0.000000   | 1.000000   |
| 75%   | 16500.000000 | 12.368421    | 0.000000   | 1.000000   |
| max   | 45400.000000 | 18.076923    | 1.000000   | 1.000000   |

The default setting of “describe” skips variables of type object. We can apply the method “describe” on the variables of type ‘object’ as follows:

```
[74]: df.describe(include=['object'])
```

```
[74]:
```

|        |        |            |              |            |                |
|--------|--------|------------|--------------|------------|----------------|
|        | make   | aspiration | num-of-doors | body-style | drive-wheels \ |
| count  | 201    | 201        | 201          | 201        | 201            |
| unique | 22     | 2          | 2            | 5          | 3              |
| top    | toyota | std        | four         | sedan      | fwd            |

|      |    |     |     |    |     |
|------|----|-----|-----|----|-----|
| freq | 32 | 165 | 115 | 94 | 118 |
|------|----|-----|-----|----|-----|

|        |                 |             |                  |             |   |
|--------|-----------------|-------------|------------------|-------------|---|
|        | engine-location | engine-type | num-of-cylinders | fuel-system | \ |
| count  | 201             | 201         | 201              | 201         |   |
| unique | 2               | 6           | 7                | 8           |   |
| top    | front           | ohc         | four             | mpfi        |   |
| freq   | 198             | 145         | 157              | 92          |   |

|        |                   |
|--------|-------------------|
|        | horsepower-binned |
| count  | 200               |
| unique | 3                 |
| top    | Low               |
| freq   | 115               |

### Value Counts

Value-counts is a good way of understanding how many units of each characteristic/variable we have. We can apply the “value\_counts” method on the column ‘drive-wheels’. Don’t forget the method “value\_counts” only works on Pandas series, not Pandas Dataframes. As a result, we only include one bracket “df[‘drive-wheels’]” not two brackets “df[[‘drive-wheels’]]”.

```
[75]: df['drive-wheels'].value_counts()
```

```
[75]: fwd      118
      rwd       75
      4wd        8
      Name: drive-wheels, dtype: int64
```

We can convert the series to a Dataframe as follows :

```
[76]: df['drive-wheels'].value_counts().to_frame()
```

```
[76]:      drive-wheels
      fwd           118
      rwd            75
      4wd             8
```

Let’s repeat the above steps but save the results to the dataframe “drive\_wheels\_counts” and rename the column ‘drive-wheels’ to ‘value\_counts’.

```
[77]: drive_wheels_counts = df['drive-wheels'].value_counts().to_frame()
      drive_wheels_counts.rename(columns={'drive-wheels': 'value_counts'},
      inplace=True)
      drive_wheels_counts
```

```
[77]:      value_counts
      fwd           118
      rwd            75
      4wd             8
```

Now let's rename the index to 'drive-wheels':

```
[78]: drive_wheels_counts.index.name = 'drive-wheels'
      drive_wheels_counts
```

```
[78]:          value_counts
drive-wheels
fwd          118
rwd           75
4wd           8
```

We can repeat the above process for the variable 'engine-location'.

```
[79]: # engine-location as variable
engine_loc_counts = df['engine-location'].value_counts().to_frame()
engine_loc_counts.rename(columns={'engine-location': 'value_counts'},
                           inplace=True)
engine_loc_counts.index.name = 'engine-location'
engine_loc_counts.head(10)
```

```
[79]:          value_counts
engine-location
front          198
rear            3
```

Examining the value counts of the engine location would not be a good predictor variable for the price. This is because we only have three cars with a rear engine and 198 with an engine in the front, this result is skewed. Thus, we are not able to draw any conclusions about the engine location.

#### 4. Basics of Grouping

The “groupby” method groups data by different categories. The data is grouped based on one or several variables and analysis is performed on the individual groups.

For example, let's group by the variable “drive-wheels”. We see that there are 3 different categories of drive wheels.

```
[80]: df['drive-wheels'].unique()
```

```
[80]: array(['rwd', 'fwd', '4wd'], dtype=object)
```

If we want to know, on average, which type of drive wheel is most valuable, we can group “drive-wheels” and then average them.

We can select the columns 'drive-wheels', 'body-style' and 'price', then assign it to the variable “df\_group\_one”.

```
[81]: df_group_one = df[['drive-wheels', 'body-style', 'price']]
```

We can then calculate the average price for each of the different categories of data.

```
[82]: # grouping results
df_group_one = df_group_one.groupby(['drive-wheels'],as_index=False).mean()
df_group_one
```

```
[82]:   drive-wheels      price
0         4wd  10241.000000
1         fwd   9244.779661
2         rwd  19757.613333
```

From our data, it seems rear-wheel drive vehicles are, on average, the most expensive, while 4-wheel and front-wheel are approximately the same in price.

You can also group with multiple variables. For example, let's group by both 'drive-wheels' and 'body-style'. This groups the dataframe by the unique combinations 'drive-wheels' and 'body-style'. We can store the results in the variable 'grouped\_test1'.

```
[83]: # grouping results
df_gptest = df[['drive-wheels','body-style','price']]
grouped_test1 = df_gptest.groupby(['drive-wheels','body-style'],as_index=False).
    ↪mean()
grouped_test1
```

```
[83]:   drive-wheels  body-style      price
0         4wd    hatchback  7603.000000
1         4wd      sedan    12647.333333
2         4wd      wagon    9095.750000
3         fwd  convertible  11595.000000
4         fwd    hardtop    8249.000000
5         fwd    hatchback   8396.387755
6         fwd      sedan    9811.800000
7         fwd      wagon    9997.333333
8         rwd  convertible  23949.600000
9         rwd    hardtop   24202.714286
10        rwd    hatchback  14337.777778
11        rwd      sedan   21711.833333
12        rwd      wagon   16994.222222
```

This grouped data is much easier to visualize when it is made into a pivot table. A pivot table is like an Excel spreadsheet, with one variable along the column and another along the row. We can convert the dataframe to a pivot table using the method "pivot" to create a pivot table from the groups.

In this case, we will leave the drive-wheel variable as the rows of the table, and pivot body-style to become the columns of the table:

```
[84]: grouped_pivot = grouped_test1.pivot(index='drive-wheels',columns='body-style')
grouped_pivot
```

```
[84]:
```

|              | price       |              |              |              |
|--------------|-------------|--------------|--------------|--------------|
| body-style   | convertible | hardtop      | hatchback    | sedan        |
| drive-wheels |             |              |              |              |
| 4wd          | NaN         | NaN          | 7603.000000  | 12647.333333 |
| fwd          | 11595.0     | 8249.000000  | 8396.387755  | 9811.800000  |
| rwd          | 23949.6     | 24202.714286 | 14337.777778 | 21711.833333 |

| body-style   | wagon        |
|--------------|--------------|
| drive-wheels |              |
| 4wd          | 9095.750000  |
| fwd          | 9997.333333  |
| rwd          | 16994.222222 |

Often, we won't have data for some of the pivot cells. We can fill these missing cells with the value 0, but any other value could potentially be used as well. It should be mentioned that missing data is quite a complex subject and is an entire course on its own.

```
[85]: grouped_pivot = grouped_pivot.fillna(0) #fill missing values with 0
grouped_pivot
```

```
[85]:
```

|              | price       |              |              |              |
|--------------|-------------|--------------|--------------|--------------|
| body-style   | convertible | hardtop      | hatchback    | sedan        |
| drive-wheels |             |              |              |              |
| 4wd          | 0.0         | 0.000000     | 7603.000000  | 12647.333333 |
| fwd          | 11595.0     | 8249.000000  | 8396.387755  | 9811.800000  |
| rwd          | 23949.6     | 24202.714286 | 14337.777778 | 21711.833333 |

| body-style   | wagon        |
|--------------|--------------|
| drive-wheels |              |
| 4wd          | 9095.750000  |
| fwd          | 9997.333333  |
| rwd          | 16994.222222 |

Question 4:

Use the “groupby” function to find the average “price” of each car based on “body-style” ?

```
[86]: df_gptest2 = df[['body-style', 'price']]
grouped_test_bodystyle = df_gptest2.groupby(['body-style'], as_index= False).
    ↪mean()
grouped_test_bodystyle
```

```
[86]:
```

|   | body-style  | price        |
|---|-------------|--------------|
| 0 | convertible | 21890.500000 |
| 1 | hardtop     | 22208.500000 |
| 2 | hatchback   | 9957.441176  |

```
3      sedan  14459.755319
4      wagon  12371.960000
```

[Double-click here for the solution.](#)


If you did not import “pyplot” let’s do it again.

```
[87]: import matplotlib.pyplot as plt
      %matplotlib inline
```

Variables: Drive Wheels and Body Style vs Price

Let’s use a heat map to visualize the relationship between Body Style vs Price.

```
[88]: #use the grouped results
      plt.pcolor(grouped_pivot, cmap='RdBu')
      plt.colorbar()
      plt.show()
```



The heatmap plots the target variable (price) proportional to colour with respect to the variables ‘drive-wheel’ and ‘body-style’ in the vertical and horizontal axis respectively. This allows us to visualize how the price is related to ‘drive-wheel’ and ‘body-style’.

The default labels convey no useful information to us. Let’s change that:

```
[89]: fig, ax = plt.subplots()
      im = ax.pcolor(grouped_pivot, cmap='RdBu')
```

```

#label names
row_labels = grouped_pivot.columns.levels[1]
col_labels = grouped_pivot.index


#move ticks and labels to the center
ax.set_xticks(np.arange(grouped_pivot.shape[1]) + 0.5, minor=False)
ax.set_yticks(np.arange(grouped_pivot.shape[0]) + 0.5, minor=False)

#insert labels
ax.set_xticklabels(row_labels, minor=False)
ax.set_yticklabels(col_labels, minor=False)

#rotate label if too long
plt.xticks(rotation=90)

fig.colorbar(im)
plt.show()

```



Visualization is very important in data science, and Python visualization packages provide great freedom. We will go more in-depth in a separate Python Visualizations course.

The main question we want to answer in this module, is “What are the main characteristics which have the most impact on the car price?”.



To get a better measure of the important characteristics, we look at the correlation of these variables with the car price, in other words: how is the car price dependent on this variable?

## 5. Correlation and Causation

Correlation: a measure of the extent of interdependence between variables.

Causation: the relationship between cause and effect between two variables.

It is important to know the difference between these two and that correlation does not imply causation. Determining correlation is much simpler than determining causation as causation may require independent experimentation.

### Pearson Correlation

The Pearson Correlation measures the linear dependence between two variables X and Y.

The resulting coefficient is a value between -1 and 1 inclusive, where:

1: Total positive linear correlation.

0: No linear correlation, the two variables most likely do not affect each other.

-1: Total negative linear correlation.

Pearson Correlation is the default method of the function “corr”. Like before we can calculate the Pearson Correlation of the of the ‘int64’ or ‘float64’ variables.

```
[90]: df.corr()
```

```
[90]:
```

|                   | symboling | normalized-losses | wheel-base | length    | \ |
|-------------------|-----------|-------------------|------------|-----------|---|
| symboling         | 1.000000  | 0.466264          | -0.535987  | -0.365404 |   |
| normalized-losses | 0.466264  | 1.000000          | -0.056661  | 0.019424  |   |
| wheel-base        | -0.535987 | -0.056661         | 1.000000   | 0.876024  |   |
| length            | -0.365404 | 0.019424          | 0.876024   | 1.000000  |   |
| width             | -0.242423 | 0.086802          | 0.814507   | 0.857170  |   |
| height            | -0.550160 | -0.373737         | 0.590742   | 0.492063  |   |
| curb-weight       | -0.233118 | 0.099404          | 0.782097   | 0.880665  |   |
| engine-size       | -0.110581 | 0.112360          | 0.572027   | 0.685025  |   |
| bore              | -0.140019 | -0.029862         | 0.493244   | 0.608971  |   |
| stroke            | -0.008245 | 0.055563          | 0.158502   | 0.124139  |   |
| compression-ratio | -0.182196 | -0.114713         | 0.250313   | 0.159733  |   |
| horsepower        | 0.075819  | 0.217299          | 0.371147   | 0.579821  |   |
| peak-rpm          | 0.279740  | 0.239543          | -0.360305  | -0.285970 |   |
| city-mpg          | -0.035527 | -0.225016         | -0.470606  | -0.665192 |   |
| highway-mpg       | 0.036233  | -0.181877         | -0.543304  | -0.698142 |   |
| price             | -0.082391 | 0.133999          | 0.584642   | 0.690628  |   |
| city-L/100km      | 0.066171  | 0.238567          | 0.476153   | 0.657373  |   |
| diesel            | -0.196735 | -0.101546         | 0.307237   | 0.211187  |   |
| gas               | 0.196735  | 0.101546          | -0.307237  | -0.211187 |   |

|           | width     | height    | curb-weight | engine-size | bore      | \ |
|-----------|-----------|-----------|-------------|-------------|-----------|---|
| symboling | -0.242423 | -0.550160 | -0.233118   | -0.110581   | -0.140019 |   |

|                   |           |           |           |           |           |
|-------------------|-----------|-----------|-----------|-----------|-----------|
| normalized-losses | 0.086802  | -0.373737 | 0.099404  | 0.112360  | -0.029862 |
| wheel-base        | 0.814507  | 0.590742  | 0.782097  | 0.572027  | 0.493244  |
| length            | 0.857170  | 0.492063  | 0.880665  | 0.685025  | 0.608971  |
| width             | 1.000000  | 0.306002  | 0.866201  | 0.729436  | 0.544885  |
| height            | 0.306002  | 1.000000  | 0.307581  | 0.074694  | 0.180449  |
| curb-weight       | 0.866201  | 0.307581  | 1.000000  | 0.849072  | 0.644060  |
| engine-size       | 0.729436  | 0.074694  | 0.849072  | 1.000000  | 0.572609  |
| bore              | 0.544885  | 0.180449  | 0.644060  | 0.572609  | 1.000000  |
| stroke            | 0.188829  | -0.062704 | 0.167562  | 0.209523  | -0.055390 |
| compression-ratio | 0.189867  | 0.259737  | 0.156433  | 0.028889  | 0.001263  |
| horsepower        | 0.615077  | -0.087027 | 0.757976  | 0.822676  | 0.566936  |
| peak-rpm          | -0.245800 | -0.309974 | -0.279361 | -0.256733 | -0.267392 |
| city-mpg          | -0.633531 | -0.049800 | -0.749543 | -0.650546 | -0.582027 |
| highway-mpg       | -0.680635 | -0.104812 | -0.794889 | -0.679571 | -0.591309 |
| price             | 0.751265  | 0.135486  | 0.834415  | 0.872335  | 0.543155  |
| city-L/100km      | 0.673363  | 0.003811  | 0.785353  | 0.745059  | 0.554610  |
| diesel            | 0.244356  | 0.281578  | 0.221046  | 0.070779  | 0.054458  |
| gas               | -0.244356 | -0.281578 | -0.221046 | -0.070779 | -0.054458 |

|                   |           |                   |            |           |   |
|-------------------|-----------|-------------------|------------|-----------|---|
|                   | stroke    | compression-ratio | horsepower | peak-rpm  | \ |
| symboling         | -0.008245 | -0.182196         | 0.075819   | 0.279740  |   |
| normalized-losses | 0.055563  | -0.114713         | 0.217299   | 0.239543  |   |
| wheel-base        | 0.158502  | 0.250313          | 0.371147   | -0.360305 |   |
| length            | 0.124139  | 0.159733          | 0.579821   | -0.285970 |   |
| width             | 0.188829  | 0.189867          | 0.615077   | -0.245800 |   |
| height            | -0.062704 | 0.259737          | -0.087027  | -0.309974 |   |
| curb-weight       | 0.167562  | 0.156433          | 0.757976   | -0.279361 |   |
| engine-size       | 0.209523  | 0.028889          | 0.822676   | -0.256733 |   |
| bore              | -0.055390 | 0.001263          | 0.566936   | -0.267392 |   |
| stroke            | 1.000000  | 0.187923          | 0.098462   | -0.065713 |   |
| compression-ratio | 0.187923  | 1.000000          | -0.214514  | -0.435780 |   |
| horsepower        | 0.098462  | -0.214514         | 1.000000   | 0.107885  |   |
| peak-rpm          | -0.065713 | -0.435780         | 0.107885   | 1.000000  |   |
| city-mpg          | -0.034696 | 0.331425          | -0.822214  | -0.115413 |   |
| highway-mpg       | -0.035201 | 0.268465          | -0.804575  | -0.058598 |   |
| price             | 0.082310  | 0.071107          | 0.809575   | -0.101616 |   |
| city-L/100km      | 0.037300  | -0.299372         | 0.889488   | 0.115830  |   |
| diesel            | 0.241303  | 0.985231          | -0.169053  | -0.475812 |   |
| gas               | -0.241303 | -0.985231         | 0.169053   | 0.475812  |   |

|                   |           |             |           |              |           |   |
|-------------------|-----------|-------------|-----------|--------------|-----------|---|
|                   | city-mpg  | highway-mpg | price     | city-L/100km | diesel    | \ |
| symboling         | -0.035527 | 0.036233    | -0.082391 | 0.066171     | -0.196735 |   |
| normalized-losses | -0.225016 | -0.181877   | 0.133999  | 0.238567     | -0.101546 |   |
| wheel-base        | -0.470606 | -0.543304   | 0.584642  | 0.476153     | 0.307237  |   |
| length            | -0.665192 | -0.698142   | 0.690628  | 0.657373     | 0.211187  |   |
| width             | -0.633531 | -0.680635   | 0.751265  | 0.673363     | 0.244356  |   |
| height            | -0.049800 | -0.104812   | 0.135486  | 0.003811     | 0.281578  |   |

|                   |           |           |           |           |           |
|-------------------|-----------|-----------|-----------|-----------|-----------|
| curb-weight       | -0.749543 | -0.794889 | 0.834415  | 0.785353  | 0.221046  |
| engine-size       | -0.650546 | -0.679571 | 0.872335  | 0.745059  | 0.070779  |
| bore              | -0.582027 | -0.591309 | 0.543155  | 0.554610  | 0.054458  |
| stroke            | -0.034696 | -0.035201 | 0.082310  | 0.037300  | 0.241303  |
| compression-ratio | 0.331425  | 0.268465  | 0.071107  | -0.299372 | 0.985231  |
| horsepower        | -0.822214 | -0.804575 | 0.809575  | 0.889488  | -0.169053 |
| peak-rpm          | -0.115413 | -0.058598 | -0.101616 | 0.115830  | -0.475812 |
| city-mpg          | 1.000000  | 0.972044  | -0.686571 | -0.949713 | 0.265676  |
| highway-mpg       | 0.972044  | 1.000000  | -0.704692 | -0.930028 | 0.198690  |
| price             | -0.686571 | -0.704692 | 1.000000  | 0.789898  | 0.110326  |
| city-L/100km      | -0.949713 | -0.930028 | 0.789898  | 1.000000  | -0.241282 |
| diesel            | 0.265676  | 0.198690  | 0.110326  | -0.241282 | 1.000000  |
| gas               | -0.265676 | -0.198690 | -0.110326 | 0.241282  | -1.000000 |

|                   |           |
|-------------------|-----------|
|                   | gas       |
| symboling         | 0.196735  |
| normalized-losses | 0.101546  |
| wheel-base        | -0.307237 |
| length            | -0.211187 |
| width             | -0.244356 |
| height            | -0.281578 |
| curb-weight       | -0.221046 |
| engine-size       | -0.070779 |
| bore              | -0.054458 |
| stroke            | -0.241303 |
| compression-ratio | -0.985231 |
| horsepower        | 0.169053  |
| peak-rpm          | 0.475812  |
| city-mpg          | -0.265676 |
| highway-mpg       | -0.198690 |
| price             | -0.110326 |
| city-L/100km      | 0.241282  |
| diesel            | -1.000000 |
| gas               | 1.000000  |

sometimes we would like to know the significant of the correlation estimate.

P-value:

What is this P-value? The P-value is the probability value that the correlation between these two variables is statistically significant. Normally, we choose a significance level of 0.05, which means that we are 95% confident that the correlation between the variables is significant.

By convention, when the

p-value is  $< 0.001$ : we say there is strong evidence that the correlation is significant.

the p-value is  $< 0.05$ : there is moderate evidence that the correlation is significant.

the p-value is  $< 0.1$ : there is weak evidence that the correlation is significant.

the p-value is  $> 0.1$ : there is no evidence that the correlation is significant.

We can obtain this information using “stats” module in the “scipy” library.

```
[91]: from scipy import stats
```

Wheel-base vs Price

Let's calculate the Pearson Correlation Coefficient and P-value of 'wheel-base' and 'price'.

```
[92]: pearson_coef, p_value = stats.pearsonr(df['wheel-base'], df['price'])
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value_
      ↳of P =", p_value)
```

The Pearson Correlation Coefficient is 0.584641822265508 with a P-value of P = 8.076488270733218e-20

Conclusion:

Since the p-value is  $< 0.001$ , the correlation between wheel-base and price is statistically significant, although the linear relationship isn't extremely strong ( $\sim 0.585$ )

Horsepower vs Price

Let's calculate the Pearson Correlation Coefficient and P-value of 'horsepower' and 'price'.

```
[93]: pearson_coef, p_value = stats.pearsonr(df['horsepower'], df['price'])
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value_
      ↳of P = ", p_value)
```

The Pearson Correlation Coefficient is 0.8095745670036559 with a P-value of P = 6.369057428260101e-48

Conclusion:

Since the p-value is  $< 0.001$ , the correlation between horsepower and price is statistically significant, and the linear relationship is quite strong ( $\sim 0.809$ , close to 1)

Length vs Price

Let's calculate the Pearson Correlation Coefficient and P-value of 'length' and 'price'.

```
[94]: pearson_coef, p_value = stats.pearsonr(df['length'], df['price'])
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value_
      ↳of P = ", p_value)
```

The Pearson Correlation Coefficient is 0.6906283804483638 with a P-value of P = 8.016477466159556e-30

Conclusion:

Since the p-value is  $< 0.001$ , the correlation between length and price is statistically significant, and the linear relationship is moderately strong ( $\sim 0.691$ ).

Width vs Price

Let's calculate the Pearson Correlation Coefficient and P-value of 'width' and 'price':

```
[95]: pearson_coef, p_value = stats.pearsonr(df['width'], df['price'])
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value_
      ↳of P =", p_value )
```

The Pearson Correlation Coefficient is 0.7512653440522673 with a P-value of P = 9.200335510481646e-38

**Conclusion:** Since the p-value is  $< 0.001$ , the correlation between width and price is statistically significant, and the linear relationship is quite strong ( $\sim 0.751$ ).

### 0.0.1 Curb-weight vs Price

Let's calculate the Pearson Correlation Coefficient and P-value of 'curb-weight' and 'price':

```
[96]: pearson_coef, p_value = stats.pearsonr(df['curb-weight'], df['price'])
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value_
      ↳of P = ", p_value)
```

The Pearson Correlation Coefficient is 0.8344145257702843 with a P-value of P = 2.189577238894065e-53

Conclusion:

Since the p-value is  $< 0.001$ , the correlation between curb-weight and price is statistically significant, and the linear relationship is quite strong ( $\sim 0.834$ ).

Engine-size vs Price

Let's calculate the Pearson Correlation Coefficient and P-value of 'engine-size' and 'price':

```
[97]: pearson_coef, p_value = stats.pearsonr(df['engine-size'], df['price'])
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value_
      ↳of P =", p_value)
```

The Pearson Correlation Coefficient is 0.8723351674455185 with a P-value of P = 9.265491622198389e-64

Conclusion:

Since the p-value is  $< 0.001$ , the correlation between engine-size and price is statistically significant, and the linear relationship is very strong ( $\sim 0.872$ ).

Bore vs Price

Let's calculate the Pearson Correlation Coefficient and P-value of 'bore' and 'price':

```
[98]: pearson_coef, p_value = stats.pearsonr(df['bore'], df['price'])
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value_
      ↳of P = ", p_value )
```

The Pearson Correlation Coefficient is 0.5431553832626602 with a P-value of P = 8.049189483935489e-17

Conclusion:

Since the p-value is  $< 0.001$ , the correlation between bore and price is statistically significant, but the linear relationship is only moderate ( $\sim 0.521$ ).

We can relate the process for each 'City-mpg' and 'Highway-mpg':

City-mpg vs Price

```
[99]: pearson_coef, p_value = stats.pearsonr(df['city-mpg'], df['price'])
      print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value_
      ↳of P = ", p_value)
```

The Pearson Correlation Coefficient is -0.6865710067844678 with a P-value of P = 2.321132065567641e-29

Conclusion:

Since the p-value is  $< 0.001$ , the correlation between city-mpg and price is statistically significant, and the coefficient of  $\sim -0.687$  shows that the relationship is negative and moderately strong.

Highway-mpg vs Price

```
[100]: pearson_coef, p_value = stats.pearsonr(df['highway-mpg'], df['price'])
      print( "The Pearson Correlation Coefficient is", pearson_coef, " with a P-value_
      ↳of P = ", p_value )
```

The Pearson Correlation Coefficient is -0.704692265058953 with a P-value of P = 1.7495471144476358e-31

**Conclusion:** Since the p-value is  $< 0.001$ , the correlation between highway-mpg and price is statistically significant, and the coefficient of  $\sim -0.705$  shows that the relationship is negative and moderately strong.

## 6. ANOVA

ANOVA: Analysis of Variance

The Analysis of Variance (ANOVA) is a statistical method used to test whether there are significant differences between the means of two or more groups. ANOVA returns two parameters:

F-test score: ANOVA assumes the means of all groups are the same, calculates how much the actual means deviate from the assumption, and reports it as the F-test score. A larger score means there is a larger difference between the means.

P-value: P-value tells how statistically significant is our calculated score value.

If our price variable is strongly correlated with the variable we are analyzing, expect ANOVA to return a sizeable F-test score and a small p-value.

Drive Wheels

Since ANOVA analyzes the difference between different groups of the same variable, the groupby function will come in handy. Because the ANOVA algorithm averages the data automatically, we do not need to take the average before hand.

Let's see if different types 'drive-wheels' impact 'price', we group the data.

Let's see if different types 'drive-wheels' impact 'price', we group the data.

```
[101]: grouped_test2=df_gpptest[['drive-wheels', 'price']].groupby(['drive-wheels'])
grouped_test2.head(2)
```

```
[101]:   drive-wheels  price
0          rwd  13495.0
1          rwd  16500.0
3          fwd  13950.0
4          4wd  17450.0
5          fwd  15250.0
136         4wd   7603.0
```

```
[102]: df_gpptest
```

```
[102]:   drive-wheels  body-style  price
0          rwd  convertible  13495.0
1          rwd  convertible  16500.0
2          rwd   hatchback  16500.0
3          fwd         sedan  13950.0
4          4wd         sedan  17450.0
..         ...           ...     ...
196         rwd         sedan  16845.0
197         rwd         sedan  19045.0
198         rwd         sedan  21485.0
199         rwd         sedan  22470.0
200         rwd         sedan  22625.0
```

```
[201 rows x 3 columns]
```

We can obtain the values of the method group using the method "get\_group".

```
[103]: grouped_test2.get_group('4wd')['price']
```

```
[103]: 4          17450.0
136         7603.0
140         9233.0
141        11259.0
144         8013.0
145        11694.0
150         7898.0
151         8778.0
Name: price, dtype: float64
```

we can use the function 'f\_oneway' in the module 'stats' to obtain the F-test score and P-value.

```
[104]: # ANOVA
f_val, p_val = stats.f_oneway(grouped_test2.get_group('fwd')['price'],
    ↳ grouped_test2.get_group('rwd')['price'], grouped_test2.
    ↳ get_group('4wd')['price'])

print( "ANOVA results: F=", f_val, ", P =", p_val)
```

ANOVA results: F= 67.95406500780399 , P = 3.3945443577151245e-23

This is a great result, with a large F test score showing a strong correlation and a P value of almost 0 implying almost certain statistical significance. But does this mean all three tested groups are all this highly correlated?

**Separately: fwd and rwd**

```
[105]: f_val, p_val = stats.f_oneway(grouped_test2.get_group('fwd')['price'],
    ↳ grouped_test2.get_group('rwd')['price'])

print( "ANOVA results: F=", f_val, ", P =", p_val )
```

ANOVA results: F= 130.5533160959111 , P = 2.2355306355677845e-23

Let's examine the other groups

**4wd and rwd**

```
[106]: f_val, p_val = stats.f_oneway(grouped_test2.get_group('4wd')['price'],
    ↳ grouped_test2.get_group('rwd')['price'])

print( "ANOVA results: F=", f_val, ", P =", p_val)
```

ANOVA results: F= 8.580681368924756 , P = 0.004411492211225333

4wd and fwd

```
[107]: f_val, p_val = stats.f_oneway(grouped_test2.get_group('4wd')['price'],
    ↳ grouped_test2.get_group('fwd')['price'])

print("ANOVA results: F=", f_val, ", P =", p_val)
```

ANOVA results: F= 0.665465750252303 , P = 0.41620116697845666

Conclusion: Important Variables

We now have a better idea of what our data looks like and which variables are important to take into account when predicting the car price. We have narrowed it down to the following variables:

Continuous numerical variables:

Length

Width



Curb-weight

Engine-size

Horsepower

City-mpg

Highway-mpg

Wheel-base

Bore

Categorical variables:

Drive-wheels

As we now move into building machine learning models to automate our analysis, feeding the model with variables that meaningfully affect our target variable will improve our model's prediction performance.

<p><a href="https://cocl.us/corsera\_da0101en\_notebook\_bottom"><img src="https://s3-api.us-geo.

About the Authors:

This notebook was written by Mahdi Noorian PhD, Joseph Santarcangelo, Bahare Talayian, Eric Xiao, Steven Dong, Parizad, Hima Vsudevan and Fiorella Wenver and Yi Yao.

Joseph Santarcangelo is a Data Scientist at IBM, and holds a PhD in Electrical Engineering. His research focused on using Machine Learning, Signal Processing, and Computer Vision to determine how videos impact human cognition. Joseph has been working for IBM since he completed his PhD.

Copyright © 2018 IBM Developer Skills Network. This notebook and its source code are released under the terms of the MIT License.

# model-evaluation-and-refinement

June 21, 2020

## Module 5: Model Evaluation and Refinement

We have built models and made predictions of vehicle prices. Now we will determine how accurate these predictions are.

This dataset was hosted on IBM Cloud object click [HERE](#) for free storage.

```
[1]: import pandas as pd
import numpy as np

# Import clean data
path = 'https://s3-api.us-geo.objectstorage.softlayer.net/cf-courses-data/
↪CognitiveClass/DA0101EN/module_5_auto.csv'
df = pd.read_csv(path)
```

```
[2]: df.to_csv('module_5_auto.csv')
```

First lets only use numeric data

```
[3]: df=df._get_numeric_data()
df.head()
```

```
[3]:
```

|   | Unnamed: 0 | Unnamed: 0.1 | symboling | normalized-losses | wheel-base | \ |
|---|------------|--------------|-----------|-------------------|------------|---|
| 0 | 0          | 0            | 3         | 122               | 88.6       |   |
| 1 | 1          | 1            | 3         | 122               | 88.6       |   |
| 2 | 2          | 2            | 1         | 122               | 94.5       |   |
| 3 | 3          | 3            | 2         | 164               | 99.8       |   |
| 4 | 4          | 4            | 2         | 164               | 99.4       |   |

|   | length   | width    | height | curb-weight | engine-size | ... | stroke | \ |
|---|----------|----------|--------|-------------|-------------|-----|--------|---|
| 0 | 0.811148 | 0.890278 | 48.8   | 2548        | 130         | ... | 2.68   |   |
| 1 | 0.811148 | 0.890278 | 48.8   | 2548        | 130         | ... | 2.68   |   |
| 2 | 0.822681 | 0.909722 | 52.4   | 2823        | 152         | ... | 3.47   |   |
| 3 | 0.848630 | 0.919444 | 54.3   | 2337        | 109         | ... | 3.40   |   |
| 4 | 0.848630 | 0.922222 | 54.3   | 2824        | 136         | ... | 3.40   |   |

|   | compression-ratio | horsepower | peak-rpm | city-mpg | highway-mpg | price   | \ |
|---|-------------------|------------|----------|----------|-------------|---------|---|
| 0 | 9.0               | 111.0      | 5000.0   | 21       | 27          | 13495.0 |   |
| 1 | 9.0               | 111.0      | 5000.0   | 21       | 27          | 16500.0 |   |

|   |      |       |        |    |    |         |
|---|------|-------|--------|----|----|---------|
| 2 | 9.0  | 154.0 | 5000.0 | 19 | 26 | 16500.0 |
| 3 | 10.0 | 102.0 | 5500.0 | 24 | 30 | 13950.0 |
| 4 | 8.0  | 115.0 | 5500.0 | 18 | 22 | 17450.0 |

|   | city-L/100km | diesel | gas |
|---|--------------|--------|-----|
| 0 | 11.190476    | 0      | 1   |
| 1 | 11.190476    | 0      | 1   |
| 2 | 12.368421    | 0      | 1   |
| 3 | 9.791667     | 0      | 1   |
| 4 | 13.055556    | 0      | 1   |

[5 rows x 21 columns]

Libraries for plotting

```
[4]: %%capture
      ! pip install ipywidgets
```

```
[5]: from IPython.display import display
      from IPython.html import widgets
      from IPython.display import display
      from ipywidgets import interact, interactive, fixed, interact_manual
```

```
/home/jupyterlab/conda/envs/python/lib/python3.6/site-
packages/IPython/html.py:14: ShimWarning: The `IPython.html` package has been
deprecated since IPython 4.0. You should import from `notebook` instead.
`IPython.html.widgets` has moved to `ipywidgets`.
  "`IPython.html.widgets` has moved to `ipywidgets`.", ShimWarning)
```

Functions for plotting

```
[6]: def DistributionPlot(RedFunction, BlueFunction, RedName, BlueName, Title):
      width = 12
      height = 10
      plt.figure(figsize=(width, height))

      ax1 = sns.distplot(RedFunction, hist=False, color="r", label=RedName)
      ax2 = sns.distplot(BlueFunction, hist=False, color="b", label=BlueName,
      ↪ax=ax1)

      plt.title(Title)
      plt.xlabel('Price (in dollars)')
      plt.ylabel('Proportion of Cars')

      plt.show()
      plt.close()
```

```
[7]: def PollyPlot(xtrain, xtest, y_train, y_test, lr, poly_transform):
    width = 12
    height = 10
    plt.figure(figsize=(width, height))

    #training data
    #testing data
    # lr: linear regression object
    #poly_transform: polynomial transformation object

    xmax=max([xtrain.values.max(), xtest.values.max()])

    xmin=min([xtrain.values.min(), xtest.values.min()])

    x=np.arange(xmin, xmax, 0.1)

    plt.plot(xtrain, y_train, 'ro', label='Training Data')
    plt.plot(xtest, y_test, 'go', label='Test Data')
    plt.plot(x, lr.predict(poly_transform.fit_transform(x.reshape(-1, 1))),
    ↪label='Predicted Function')
    plt.ylim([-10000, 60000])
    plt.ylabel('Price')
    plt.legend()
```

## Part 1: Training and Testing

An important step in testing your model is to split your data into training and testing data. We will place the target data price in a separate dataframe y:

```
[8]: y_data = df['price']
```

drop price data in x data

```
[9]: x_data=df.drop('price',axis=1)
```

Now we randomly split our data into training and testing data using the function train\_test\_split.

```
[10]: from sklearn.model_selection import train_test_split

x_train, x_test, y_train, y_test = train_test_split(x_data, y_data, test_size=0.
    ↪15, random_state=1)

print("number of test samples :", x_test.shape[0])
print("number of training samples:",x_train.shape[0])
```

```
number of test samples : 31
number of training samples: 170
```

The test\_size parameter sets the proportion of data that is split into the testing set. In the above, the testing set is set to 10% of the total dataset.

Question #1):

Use the function “train\_test\_split” to split up the data set such that 40% of the data samples will be utilized for testing, set the parameter “random\_state” equal to zero. The output of the function should be the following: “x\_train\_1” , “x\_test\_1”, “y\_train\_1” and “y\_test\_1”.

```
[11]: x_train1, x_test1, y_train1, y_test1 = train_test_split(x_data, y_data,
    ↳ test_size=0.4, random_state=0)
print("number of test samples :", x_test1.shape[0])
print("number of training samples:", x_train1.shape[0])
```

```
number of test samples : 81
number of training samples: 120
```

Double-click here for the solution.

Let's import LinearRegression from the module linear\_model.

```
[12]: from sklearn.linear_model import LinearRegression
```

We create a Linear Regression object:

```
[13]: lre=LinearRegression()
```

we fit the model using the feature horsepower

```
[14]: lre.fit(x_train[['horsepower']], y_train)
```

```
[14]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
    normalize=False)
```

Let's Calculate the  $R^2$  on the test data:

```
[15]: lre.score(x_test[['horsepower']], y_test)
```

```
[15]: 0.707688374146705
```

we can see the  $R^2$  is much smaller using the test data.

```
[16]: lre.score(x_train[['horsepower']], y_train)
```

```
[16]: 0.6449517437659684
```

Question #2):

Find the  $R^2$  on the test data using 90% of the data for training data

```
[17]: x_train1, x_test1, y_train1, y_test1 = train_test_split(x_data, y_data,
    ↪test_size=0.1, random_state=0)
lre.fit(x_train1[['horsepower']],y_train1)
lre.score(x_test1[['horsepower']],y_test1)
```

```
[17]: 0.7340722810055448
```

Double-click here for the solution.

Sometimes you do not have sufficient testing data; as a result, you may want to perform Cross-validation. Let's go over several methods that you can use for Cross-validation.

Cross-validation Score

Lets import model\_selection from the module cross\_val\_score.

```
[18]: from sklearn.model_selection import cross_val_score
```

We input the object, the feature in this case 'horsepower', the target data (y\_data). The parameter 'cv' determines the number of folds; in this case 4.

```
[19]: Rcross = cross_val_score(lre, x_data[['horsepower']], y_data, cv=4)
```

The default scoring is  $R^2$ ; each element in the array has the average  $R^2$  value in the fold:

```
[20]: Rcross
```

```
[20]: array([0.7746232 , 0.51716687, 0.74785353, 0.04839605])
```

We can calculate the average and standard deviation of our estimate:

```
[21]: print("The mean of the folds are", Rcross.mean(), "and the standard deviation,
    ↪is" , Rcross.std())
```

The mean of the folds are 0.522009915042119 and the standard deviation is 0.291183944475603

We can use negative squared error as a score by setting the parameter 'scoring' metric to 'neg\_mean\_squared\_error'.

```
[22]: -1 * cross_val_score(lre,x_data[['horsepower']],
    ↪y_data,cv=4,scoring='neg_mean_squared_error')
```

```
[22]: array([20254142.84026702, 43745493.2650517 , 12539630.34014931,
    17561927.72247591])
```

Question #3):

Calculate the average  $R^2$  using two folds, find the average  $R^2$  for the second fold utilizing the horsepower as a feature :

```
[23]: Rc=cross_val_score(lre,x_data[['horsepower']], y_data,cv=2)
      Rc.mean()
```

```
[23]: 0.5166761697127429
```

Double-click here for the solution.

You can also use the function 'cross\_val\_predict' to predict the output. The function splits up the data into the specified number of folds, using one fold to get a prediction while the rest of the folds are used as test data. First import the function:

```
[24]: from sklearn.model_selection import cross_val_predict
```

We input the object, the feature in this case 'horsepower', the target data y\_data. The parameter 'cv' determines the number of folds; in this case 4. We can produce an output:

```
[25]: yhat = cross_val_predict(lre,x_data[['horsepower']], y_data,cv=4)
      yhat[0:5]
```

```
[25]: array([14141.63807508, 14141.63807508, 20814.29423473, 12745.03562306,
           14762.35027598])
```

## Part 2: Overfitting, Underfitting and Model Selection

It turns out that the test data sometimes referred to as the out of sample data is a much better measure of how well your model performs in the real world. One reason for this is overfitting; let's go over some examples. It turns out these differences are more apparent in Multiple Linear Regression and Polynomial Regression so we will explore overfitting in that context.

Let's create Multiple linear regression objects and train the model using 'horsepower', 'curb-weight', 'engine-size' and 'highway-mpg' as features.

```
[26]: lr = LinearRegression()
      lr.fit(x_train[['horsepower', 'curb-weight', 'engine-size', 'highway-mpg']],
            y_train)
```

```
[26]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
           normalize=False)
```

Prediction using training data:

```
[27]: yhat_train = lr.predict(x_train[['horsepower', 'curb-weight', 'engine-size',
            'highway-mpg']])
      yhat_train[0:5]
```

```
[27]: array([11927.70699817, 11236.71672034, 6436.91775515, 21890.22064982,
           16667.18254832])
```

Prediction using test data:

```
[28]: yhat_test = lr.predict(x_test[['horsepower', 'curb-weight', 'engine-size', '
↳ 'highway-mpg']])
yhat_test[0:5]
```

```
[28]: array([11349.16502418,  5914.48335385, 11243.76325987,  6662.03197043,
        15555.76936275])
```

Let's perform some model evaluation using our training and testing data separately. First we import the seaborn and matplotlib library for plotting.

```
[29]: import matplotlib.pyplot as plt
      %matplotlib inline
      import seaborn as sns
```

Let's examine the distribution of the predicted values of the training data.

```
[30]: Title = 'Distribution Plot of Predicted Value Using Training Data vs Training
↳ Data Distribution'
DistributionPlot(y_train, yhat_train, "Actual Values (Train)", "Predicted
↳ Values (Train)", Title)
```






Figure 1: Plot of predicted values using the training data compared to the training data.

So far the model seems to be doing well in learning from the training dataset. But what happens when the model encounters new data from the testing dataset? When the model generates new values from the test data, we see the distribution of the predicted values is much different from the actual target values.

```
[31]: Title='Distribution Plot of Predicted Value Using Test Data vs Data_
      ↳Distribution of Test Data'
      DistributionPlot(y_test,yhat_test,"Actual Values (Test)","Predicted Values_
      ↳(Test)",Title)
```




Figure 2: Plot of predicted value using the test data compared to the test data.

Comparing Figure 1 and Figure 2; it is evident the distribution of the test data in Figure 1 is much better at fitting the data. This difference in Figure 2 is apparent where the ranges are from 5000 to 15 000. This is where the distribution shape is exceptionally different. Let's see if polynomial regression also exhibits a drop in the prediction accuracy when analysing the test dataset.

```
[32]: from sklearn.preprocessing import PolynomialFeatures
```

### Overfitting

Overfitting occurs when the model fits the noise, not the underlying process. Therefore when testing your model using the test-set, your model does not perform as well as it is modelling noise, not the underlying process that generated the relationship. Let's create a degree 5 polynomial model.

Let's use 55 percent of the data for testing and the rest for training:

```
[33]: x_train, x_test, y_train, y_test = train_test_split(x_data, y_data, test_size=0.
↪45, random_state=0)
```

We will perform a degree 5 polynomial transformation on the feature 'horse power'.

```
[35]: pr = PolynomialFeatures(degree=5)
x_train_pr = pr.fit_transform(x_train[['horsepower']])
x_test_pr = pr.fit_transform(x_test[['horsepower']])
pr
```

```
[35]: PolynomialFeatures(degree=5, include_bias=True, interaction_only=False)
```

Now let's create a linear regression model "poly" and train it.

```
[36]: poly = LinearRegression()
poly.fit(x_train_pr, y_train)
```

```
[36]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
normalize=False)
```

We can see the output of our model using the method "predict." then assign the values to "yhat".

```
[37]: yhat = poly.predict(x_test_pr)
yhat[0:5]
```

```
[37]: array([ 6728.65561887,  7307.98782321, 12213.78770965, 18893.24804015,
19995.95195136])
```

Let's take the first five predicted values and compare it to the actual targets.

```
[38]: print("Predicted values:", yhat[0:4])
print("True values:", y_test[0:4].values)
```

```
Predicted values: [ 6728.65561887  7307.98782321 12213.78770965 18893.24804015]
True values: [ 6295. 10698. 13860. 13499.]
```

We will use the function "PollyPlot" that we defined at the beginning of the lab to display the training data, testing data, and the predicted function.

```
[39]: PollyPlot(x_train[['horsepower']], x_test[['horsepower']], y_train, y_test,
↪poly, pr)
```

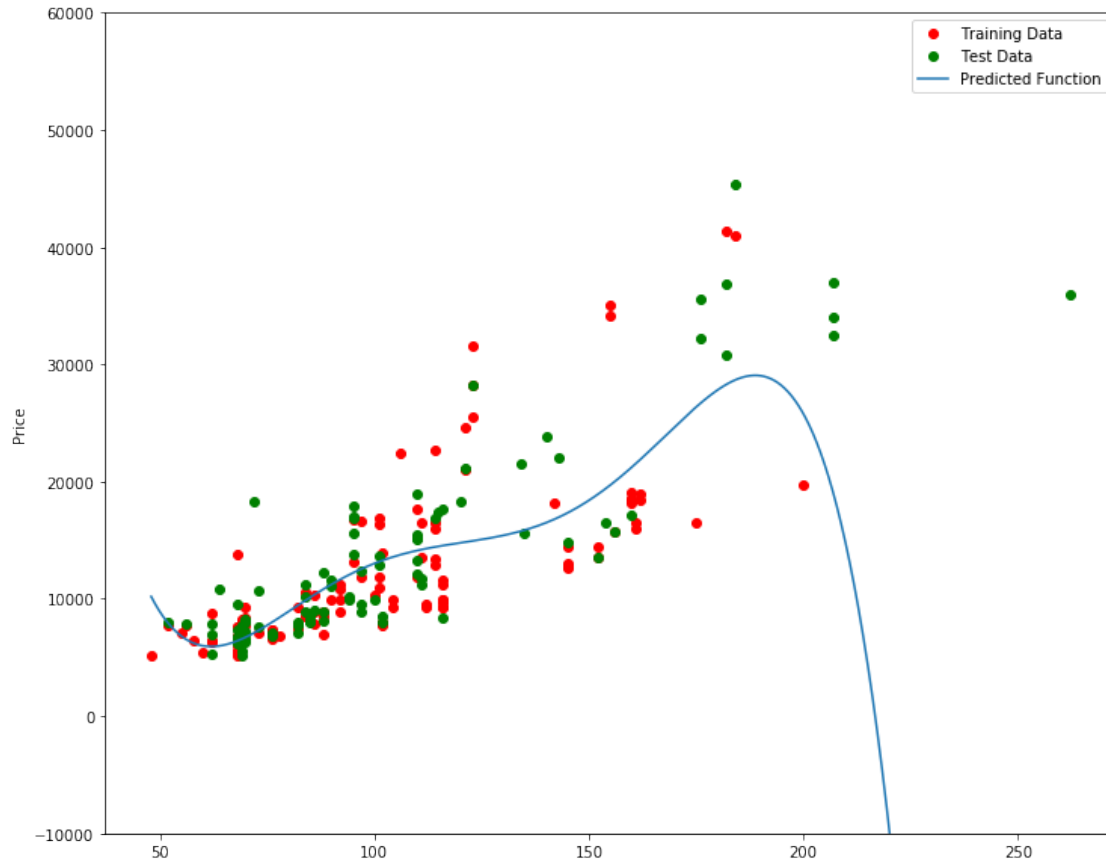


Figure 4 A polynomial regression model, red dots represent training data, green dots represent test data, and the blue line represents the model prediction.

We see that the estimated function appears to track the data but around 200 horsepower, the function begins to diverge from the data points.

$R^2$  of the training data:

```
[40]: poly.score(x_train_pr, y_train)
```

```
[40]: 0.556771690212023
```

$R^2$  of the test data:

```
[41]: poly.score(x_test_pr, y_test)
```

```
[41]: -29.871340302044153
```

We see the  $R^2$  for the training data is 0.5567 while the  $R^2$  on the test data was -29.87. The lower the  $R^2$ , the worse the model, a Negative  $R^2$  is a sign of overfitting.

Let's see how the  $R^2$  changes on the test data for different order polynomials and plot the results:

```
[42]: Rsqu_test = []

order = [1, 2, 3, 4]
for n in order:
    pr = PolynomialFeatures(degree=n)

    x_train_pr = pr.fit_transform(x_train[['horsepower']])

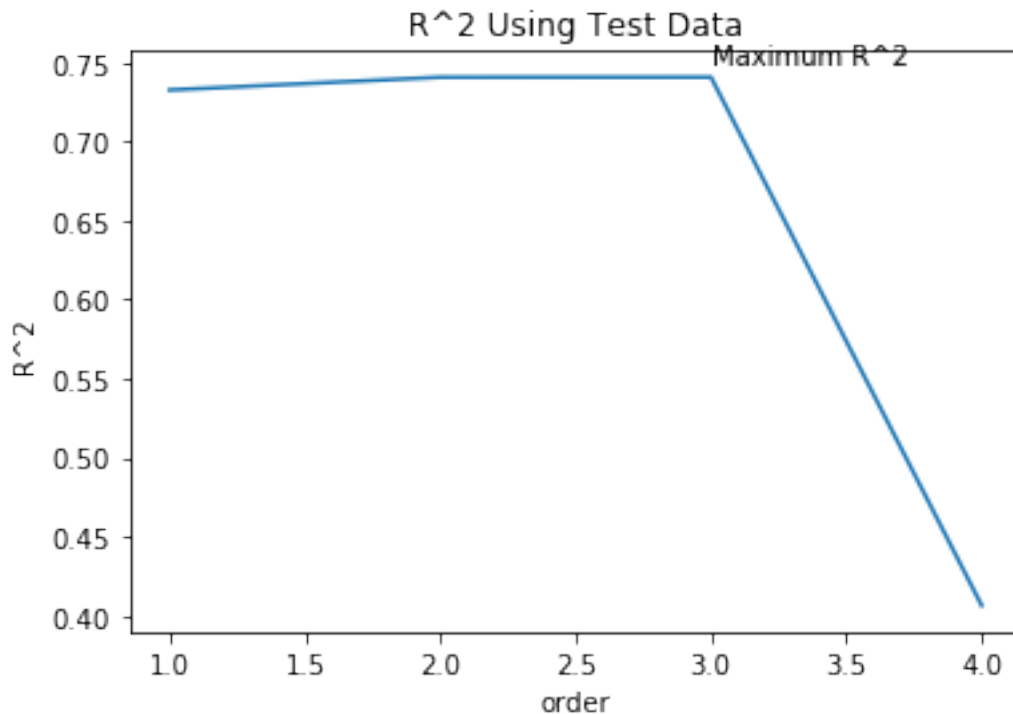
    x_test_pr = pr.fit_transform(x_test[['horsepower']])

    lr.fit(x_train_pr, y_train)

    Rsqu_test.append(lr.score(x_test_pr, y_test))

plt.plot(order, Rsqu_test)
plt.xlabel('order')
plt.ylabel('R^2')
plt.title('R^2 Using Test Data')
plt.text(3, 0.75, 'Maximum R^2 ')
```

```
[42]: Text(3, 0.75, 'Maximum R^2 ')
```



We see the  $R^2$  gradually increases until an order three polynomial is used. Then the  $R^2$  dramatically decreases at four.

The following function will be used in the next section; please run the cell.

```
[43]: def f(order, test_data):
      x_train, x_test, y_train, y_test = train_test_split(x_data, y_data,
      ↳test_size=test_data, random_state=0)
      pr = PolynomialFeatures(degree=order)
      x_train_pr = pr.fit_transform(x_train[['horsepower']])
      x_test_pr = pr.fit_transform(x_test[['horsepower']])
      poly = LinearRegression()
      poly.fit(x_train_pr, y_train)
      PollyPlot(x_train[['horsepower']], x_test[['horsepower']], y_train, y_test,
      ↳poly, pr)
```

The following interface allows you to experiment with different polynomial orders and different amounts of data.

```
[44]: interact(f, order=(0, 6, 1), test_data=(0.05, 0.95, 0.05))
```

```
interactive(children=(IntSlider(value=3, description='order', max=6), FloatSlider(value=0.45, c
```

```
[44]: <function __main__.f(order, test_data)>
```

Question #4a):

We can perform polynomial transformations with more than one feature. Create a “PolynomialFeatures” object “pr1” of degree two?

[Double-click here for the solution.](#)

Question #4b):

Transform the training and testing samples for the features ‘horsepower’, ‘curb-weight’, ‘engine-size’ and ‘highway-mpg’. Hint: use the method “fit\_transform” ?

[Double-click here for the solution.](#)

Question #4c):

How many dimensions does the new feature have? Hint: use the attribute “shape”

[Double-click here for the solution.](#)

Question #4d):

Create a linear regression model “poly1” and train the object using the method “fit” using the polynomial features?

[Double-click here for the solution.](#)

Question #4e):

Use the method “predict” to predict an output on the polynomial features, then use the function “DistributionPlot” to display the distribution of the predicted output vs the test data?

[Double-click here for the solution.](#)

Question #4f):

Use the distribution plot to determine the two regions where the predicted prices are less accurate than the actual prices.

Double-click here for the solution.

In this section, we will review Ridge Regression we will see how the parameter Alfa changes the model. Just a note here our test data will be used as validation data.

Let's perform a degree two polynomial transformation on our data.

```
[45]: pr=PolynomialFeatures(degree=2)
x_train_pr=pr.fit_transform(x_train[['horsepower', 'curb-weight', '
    ↳ 'engine-size', 'highway-mpg', 'normalized-losses', 'symboling']])
x_test_pr=pr.fit_transform(x_test[['horsepower', 'curb-weight', 'engine-size', '
    ↳ 'highway-mpg', 'normalized-losses', 'symboling']])
```

Let's import Ridge from the module linear models.

```
[46]: from sklearn.linear_model import Ridge
```

Let's create a Ridge regression object, setting the regularization parameter to 0.1

```
[47]: RigeModel=Ridge(alpha=0.1)
```

Like regular regression, you can fit the model using the method fit.

```
[49]: RigeModel.fit(x_train_pr, y_train)
```

```
/home/jupyterlab/conda/envs/python/lib/python3.6/site-
packages/sklearn/linear_model/ridge.py:125: LinAlgWarning: Ill-conditioned
matrix (rcond=1.02972e-16): result may not be accurate.
  overwrite_a=True).T
```

```
[49]: Ridge(alpha=0.1, copy_X=True, fit_intercept=True, max_iter=None,
    normalize=False, random_state=None, solver='auto', tol=0.001)
```

Similarly, you can obtain a prediction:

```
[50]: yhat = RigeModel.predict(x_test_pr)
```

Let's compare the first five predicted samples to our test set

```
[51]: print('predicted:', yhat[0:4])
print('test set :', y_test[0:4].values)
```

```
predicted: [ 6567.83081933  9597.97151399 20836.22326843 19347.69543463]
test set : [ 6295. 10698. 13860. 13499.]
```

We select the value of Alfa that minimizes the test error, for example, we can use a for loop.

```
[52]: Rsqu_test = []
Rsqu_train = []
dummy1 = []
ALFA = 10 * np.array(range(0,1000))
for alfa in ALFA:
    RigeModel = Ridge(alpha=alfa)
    RigeModel.fit(x_train_pr, y_train)
    Rsqu_test.append(RigeModel.score(x_test_pr, y_test))
    Rsqu_train.append(RigeModel.score(x_train_pr, y_train))
```

We can plot out the value of  $R^2$  for different Alphas

```
[53]: width = 12
height = 10
plt.figure(figsize=(width, height))

plt.plot(ALFA, Rsqu_test, label='validation data ')
plt.plot(ALFA, Rsqu_train, 'r', label='training Data ')
plt.xlabel('alpha')
plt.ylabel('R^2')
plt.legend()
```

```
[53]: <matplotlib.legend.Legend at 0x7f0978484630>
```




Figure 6: The blue line represents the  $R^2$  of the test data, and the red line represents the  $R^2$  of the training data. The x-axis represents the different values of Alpha

The red line in figure 6 represents the  $R^2$  of the test data, as Alpha increases the  $R^2$  decreases; therefore as Alpha increases the model performs worse on the test data. The blue line represents the  $R^2$  on the validation data, as the value for Alpha increases the  $R^2$  decreases.

Question #5):

Perform Ridge regression and calculate the  $R^2$  using the polynomial features, use the training data to train the model and test data to test the model. The parameter alpha should be set to 10.

```
[54]: RigeModel = Ridge(alpha=0)
      RigeModel.fit(x_train_pr, y_train)
      RigeModel.score(x_test_pr, y_test)
```

[54]: 0.47098333063511094

[Double-click here for the solution.](#)

Part 4: Grid Search



The term Alfa is a hyperparameter, sklearn has the class GridSearchCV to make the process of finding the best hyperparameter simpler.

Let's import GridSearchCV from the module model\_selection.

```
[55]: from sklearn.model_selection import GridSearchCV
```

We create a dictionary of parameter values:

```
[56]: parameters1= [{'alpha': [0.001,0.1,1, 10, 100, 1000, 10000, 100000, 100000]}]
parameters1
```

```
[56]: [{'alpha': [0.001, 0.1, 1, 10, 100, 1000, 10000, 100000, 100000]}]
```

Create a ridge regions object:

```
[57]: RR=Ridge()
RR
```

```
[57]: Ridge(alpha=1.0, copy_X=True, fit_intercept=True, max_iter=None,
          normalize=False, random_state=None, solver='auto', tol=0.001)
```

Create a ridge grid search object

```
[58]: Grid1 = GridSearchCV(RR, parameters1,cv=4)
```

Fit the model

```
[59]: Grid1.fit(x_data[['horsepower', 'curb-weight', 'engine-size', 'highway-mpg']],
               y_data)
```

```
/home/jupyterlab/conda/envs/python/lib/python3.6/site-
packages/sklearn/model_selection/_search.py:841: DeprecationWarning: The default
of the `iid` parameter will change from True to False in version 0.22 and will
be removed in 0.24. This will change numeric results when test-set sizes are
unequal.
```

```
DeprecationWarning)
```

```
[59]: GridSearchCV(cv=4, error_score='raise-deprecating',
                  estimator=Ridge(alpha=1.0, copy_X=True, fit_intercept=True,
max_iter=None,
normalize=False, random_state=None, solver='auto', tol=0.001),
                  fit_params=None, iid='warn', n_jobs=None,
                  param_grid=[{'alpha': [0.001, 0.1, 1, 10, 100, 1000, 10000, 100000,
100000]}],
                  pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
                  scoring=None, verbose=0)
```

The object finds the best parameter values on the validation data. We can obtain the estimator with the best parameters and assign it to the variable BestRR as follows:

```
[60]: BestRR=Grid1.best_estimator_  
BestRR
```

```
[60]: Ridge(alpha=10000, copy_X=True, fit_intercept=True, max_iter=None,  
          normalize=False, random_state=None, solver='auto', tol=0.001)
```

We now test our model on the test data

```
[61]: BestRR.score(x_test[['horsepower', 'curb-weight', 'engine-size',  
    ↪ 'highway-mpg']], y_test)
```

```
[61]: 0.8411649831036152
```

Question #6):

Perform a grid search for the alpha parameter and the normalization parameter, then find the best values of the parameters

```
[62]: parameters2= [{'alpha': [0.001,0.1,1, 10, 100,  
    ↪ 1000,10000,100000,100000], 'normalize':[True,False]} ]  
Grid2 = GridSearchCV(Ridge(), parameters2,cv=4)  
Grid2.fit(x_data[['horsepower', 'curb-weight', 'engine-size',  
    ↪ 'highway-mpg']],y_data)  
Grid2.best_estimator_
```

/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/sklearn/model\_selection/\_search.py:841: DeprecationWarning: The default of the `iid` parameter will change from True to False in version 0.22 and will be removed in 0.24. This will change numeric results when test-set sizes are unequal.

DeprecationWarning)

```
[62]: Ridge(alpha=0.1, copy_X=True, fit_intercept=True, max_iter=None,  
          normalize=True, random_state=None, solver='auto', tol=0.001)
```

Double-click here for the solution.

<p><a href="https://cocl.us/corsera\_da0101en\_notebook\_bottom"><img src="https://s3-api.us-geo.

About the Authors:

This notebook was written by Mahdi Noorian PhD, Joseph Santarcangelo, Bahare Talayian, Eric Xiao, Steven Dong, Parizad, Hima Vsudevan and Fiorella Wenver and Yi Yao.

Joseph Santarcangelo is a Data Scientist at IBM, and holds a PhD in Electrical Engineering. His research focused on using Machine Learning, Signal Processing, and Computer Vision to determine how videos impact human cognition. Joseph has been working for IBM since he completed his PhD.

Copyright © 2018 IBM Developer Skills Network. This notebook and its source code are released under the terms of the MIT License.

# model-development

June 21, 2020

## Module 4: Model Development

In this section, we will develop several models that will predict the price of the car using the variables or features. This is just an estimate but should give us an objective idea of how much the car should cost.

Some questions we want to ask in this module

do I know if the dealer is offering fair value for my trade-in?

do I know if I put a fair value on my car?

Data Analytics, we often use Model Development to help us predict future observations from the data we have.

A Model will help us understand the exact relationship between different variables and how these variables are used to predict the result.

Setup

Import libraries

```
[3]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

load data and store in dataframe df:

This dataset was hosted on IBM Cloud object click [HERE](#) for free storage.

```
[4]: # path of data
path = 'https://s3-api.us-geo.objectstorage.softlayer.net/cf-courses-data/
↪CognitiveClass/DA0101EN/automobileEDA.csv'
df = pd.read_csv(path)
df.head()
```

```
[4]:   symboling  normalized-losses   make aspiration num-of-doors \
0         3         122  alfa-romero      std         two
1         3         122  alfa-romero      std         two
2         1         122  alfa-romero      std         two
3         2         164      audi      std         four
4         2         164      audi      std         four
```

|   | body-style  | drive-wheels | engine-location | wheel-base | length   | ... | \ |
|---|-------------|--------------|-----------------|------------|----------|-----|---|
| 0 | convertible | rwd          | front           | 88.6       | 0.811148 | ... |   |
| 1 | convertible | rwd          | front           | 88.6       | 0.811148 | ... |   |
| 2 | hatchback   | rwd          | front           | 94.5       | 0.822681 | ... |   |
| 3 | sedan       | fwd          | front           | 99.8       | 0.848630 | ... |   |
| 4 | sedan       | 4wd          | front           | 99.4       | 0.848630 | ... |   |

|   | compression-ratio | horsepower | peak-rpm | city-mpg | highway-mpg | price   | \ |
|---|-------------------|------------|----------|----------|-------------|---------|---|
| 0 | 9.0               | 111.0      | 5000.0   | 21       | 27          | 13495.0 |   |
| 1 | 9.0               | 111.0      | 5000.0   | 21       | 27          | 16500.0 |   |
| 2 | 9.0               | 154.0      | 5000.0   | 19       | 26          | 16500.0 |   |
| 3 | 10.0              | 102.0      | 5500.0   | 24       | 30          | 13950.0 |   |
| 4 | 8.0               | 115.0      | 5500.0   | 18       | 22          | 17450.0 |   |

|   | city-L/100km | horsepower-binned | diesel | gas |
|---|--------------|-------------------|--------|-----|
| 0 | 11.190476    | Medium            | 0      | 1   |
| 1 | 11.190476    | Medium            | 0      | 1   |
| 2 | 12.368421    | Medium            | 0      | 1   |
| 3 | 9.791667     | Medium            | 0      | 1   |
| 4 | 13.055556    | Medium            | 0      | 1   |

[5 rows x 29 columns]

## 1. Linear Regression and Multiple Linear Regression

### Linear Regression

One example of a Data Model that we will be using is

Simple Linear Regression.

Simple Linear Regression is a method to help us understand the relationship between two variables:

The predictor/independent variable (X)

The response/dependent variable (that we want to predict)(Y)

The result of Linear Regression is a linear function that predicts the response (dependent) variable as a function of the predictor (independent) variable.

$$Y : \text{Response Variable} \quad X : \text{Predictor Variables}$$

Linear function:

$$\hat{Y} = a + bX$$

a refers to the intercept of the regression line, in other words: the value of Y when X is 0

b refers to the slope of the regression line, in other words: the value with which Y changes when X increases by 1 unit

Lets load the modules for linear regression

```
[8]: from sklearn.linear_model import LinearRegression
```

Create the linear regression object

```
[7]: lm = LinearRegression()  
lm
```

```
[7]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,  
normalize=False)
```

How could Highway-mpg help us predict car price?

For this example, we want to look at how highway-mpg can help us predict car price. Using simple linear regression, we will create a linear function with “highway-mpg” as the predictor variable and the “price” as the response variable.

```
[9]: X = df[['highway-mpg']]  
Y = df['price']
```

Fit the linear model using highway-mpg.

```
[10]: lm.fit(X,Y)
```

```
[10]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,  
normalize=False)
```

We can output a prediction

```
[11]: Yhat=lm.predict(X)  
Yhat[0:5]
```

```
[11]: array([16236.50464347, 16236.50464347, 17058.23802179, 13771.3045085 ,  
20345.17153508])
```

What is the value of the intercept (a)?

```
[12]: lm.intercept_
```

```
[12]: 38423.3058581574
```

What is the value of the Slope (b)?

```
[13]: lm.coef_
```

```
[13]: array([-821.73337832])
```

What is the final estimated linear model we get?

As we saw above, we should get a final linear model with the structure:

$$\hat{Y} = a + bX$$

Plugging in the actual values we get:

price = 38423.31 - 821.73 x highway-mpg

Question #1 a):

Create a linear regression object?

```
[14]: lm1 = LinearRegression()
      lm1
```

```
[14]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
      normalize=False)
```

Double-click here for the solution.

Question #1 b):

Train the model using 'engine-size' as the independent variable and 'price' as the dependent variable?

```
[22]: x = df[["engine-size"]]
      y = df[["price"]]
      lm1.fit(x,y)
```

```
[22]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
      normalize=False)
```

Double-click here for the solution.

Question #1 c):

Find the slope and intercept of the model?

Slope

```
[27]: yhat = lm1.predict(x)
      lm1.coef_
```

```
[27]: array([[166.86001569]])
```

Intercept

```
[28]: # Write your code below and press Shift+Enter to execute
      lm1.intercept_
```

```
[28]: array([-7963.33890628])
```

Double-click here for the solution.

Question #1 d):

What is the equation of the predicted line. You can use  $x$  and  $\hat{y}$  or ‘engine-size’ or ‘price’?

## 1 You can type you answer here

Double-click here for the solution.

Multiple Linear Regression

What if we want to predict car price using more than one variable?

If we want to use more variables in our model to predict car price, we can use Multiple Linear Regression. Multiple Linear Regression is very similar to Simple Linear Regression, but this method is used to explain the relationship between one continuous response (dependent) variable and two or more predictor (independent) variables. Most of the real-world regression models involve multiple predictors. We will illustrate the structure by using four predictor variables, but these results can generalize to any integer:

$Y$  : Response Variable  $X_1$  : Predictor Variable 1  $X_2$  : Predictor Variable 2  $X_3$  : Predictor Variable 3  $X_4$  : Predictor Variable 4

$a$  : intercept  $b_1$  : coefficients of Variable 1  $b_2$  : coefficients of Variable 2  $b_3$  : coefficients of Variable 3  $b_4$  : coefficients of Variable 4

The equation is given by

$$\hat{Y} = a + b_1X_1 + b_2X_2 + b_3X_3 + b_4X_4$$

From the previous section we know that other good predictors of price could be:

Horsepower

Curb-weight

Engine-size

Highway-mpg

Let’s develop a model using these variables as the predictor variables.

```
[29]: Z = df[['horsepower', 'curb-weight', 'engine-size', 'highway-mpg']]
```

Fit the linear model using the four above-mentioned variables.

```
[30]: lm.fit(Z, df['price'])
```

```
[30]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
                        normalize=False)
```

What is the value of the intercept(a)?

```
[31]: lm.intercept_
```

```
[31]: -15806.62462632922
```

What are the values of the coefficients (b1, b2, b3, b4)?

```
[33]: lm.coef_
```

```
[33]: array([53.49574423,  4.70770099, 81.53026382, 36.05748882])
```

What is the final estimated linear model that we get?

As we saw above, we should get a final linear function with the structure:

$$\hat{Y} = a + b_1X_1 + b_2X_2 + b_3X_3 + b_4X_4$$

What is the linear function we get in this example?

Price = -15678.742628061467 + 52.65851272 x horsepower + 4.69878948 x curb-weight + 81.95906216 x engine-size + 33.58258185 x highway-mpg

Question #2 a):

Create and train a Multiple Linear Regression model “lm2” where the response variable is price, and the predictor variable is ‘normalized-losses’ and ‘highway-mpg’.

```
[34]: lm2 = LinearRegression()  
lm2.fit(df[['normalized-losses' , 'highway-mpg']],df['price'])
```

```
[34]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,  
normalize=False)
```

Double-click here for the solution.

Question #2 b):

Find the coefficient of the model?

```
[36]: lm2.coef_
```

```
[36]: array([ 1.49789586, -820.45434016])
```

Double-click here for the solution.

## 2) Model Evaluation using Visualization

Now that we’ve developed some models, how do we evaluate our models and how do we choose the best one? One way to do this is by using visualization.

import the visualization package: seaborn

```
[37]: # import the visualization package: seaborn  
import seaborn as sns  
%matplotlib inline
```



## Regression Plot


When it comes to simple linear regression, an excellent way to visualize the fit of our model is by using regression plots.

This plot will show a combination of a scattered data points (a scatter plot), as well as the fitted linear regression line going through the data. This will give us a reasonable estimate of the relationship between the two variables, the strength of the correlation, as well as the direction (positive or negative correlation).

Let's visualize Horsepower as potential predictor variable of price:

```
[38]: width = 12  
      height = 10  
      plt.figure(figsize=(width, height))  
      sns.regplot(x="highway-mpg", y="price", data=df)  
      plt.ylim(0,)
```

```
[38]: (0, 48290.29917233836)
```




We can see from this plot that price is negatively correlated to highway-mpg, since the regression

slope is negative. One thing to keep in mind when looking at a regression plot is to pay attention to how scattered the data points are around the regression line. This will give you a good indication of the variance of the data, and whether a linear model would be the best fit or not. If the data is too far off from the line, this linear model might not be the best model for this data. Let's compare this plot to the regression plot of "peak-rpm".

```
[39]: plt.figure(figsize=(width, height))
      sns.regplot(x="peak-rpm", y="price", data=df)
      plt.ylim(0,)
```

```
[39]: (0, 47422.919330307624)
```



Comparing the regression plot of "peak-rpm" and "highway-mpg" we see that the points for "highway-mpg" are much closer to the generated line and on the average decrease. The points for "peak-rpm" have more spread around the predicted line, and it is much harder to determine if the points are decreasing or increasing as the "highway-mpg" increases.

Question #3:

Given the regression plots above is "peak-rpm" or "highway-mpg" more strongly correlated with "price". Use the method ".corr()" to verify your answer.

```
[40]: df[["peak-rpm", "highway-mpg", "price"]].corr()
```

```
[40]:
```

|             | peak-rpm  | highway-mpg | price     |
|-------------|-----------|-------------|-----------|
| peak-rpm    | 1.000000  | -0.058598   | -0.101616 |
| highway-mpg | -0.058598 | 1.000000    | -0.704692 |
| price       | -0.101616 | -0.704692   | 1.000000  |

Double-click here for the solution.

## Residual Plot

A good way to visualize the variance of the data is to use a residual plot.

What is a residual?

The difference between the observed value ( $y$ ) and the predicted value ( $\hat{y}$ ) is called the residual ( $e$ ). When we look at a regression plot, the residual is the distance from the data point to the fitted regression line.

So what is a residual plot?


A residual plot is a graph that shows the residuals on the vertical y-axis and the independent variable on the horizontal x-axis.

What do we pay attention to when looking at a residual plot?

We look at the spread of the residuals:

- If the points in a residual plot are randomly spread out around the x-axis, then a linear model is appropriate for the data. Why is that? Randomly spread out residuals means that the variance is constant, and thus the linear model is a good fit for this data.

```
[41]: width = 12
height = 10
plt.figure(figsize=(width, height))
sns.residplot(df['highway-mpg'], df['price'])
plt.show()
```



What is this plot telling us?

We can see from this residual plot that the residuals are not randomly spread around the x-axis, which leads us to believe that maybe a non-linear model is more appropriate for this data.

### Multiple Linear Regression

How do we visualize a model for Multiple Linear Regression? This gets a bit more complicated because you can't visualize it with regression or residual plot.

One way to look at the fit of the model is by looking at the distribution plot: We can look at the distribution of the fitted values that result from the model and compare it to the distribution of the actual values.

First lets make a prediction

```
[42]: Y_hat = lm.predict(Z)
```

```
[43]: plt.figure(figsize=(width, height))
```

```
ax1 = sns.distplot(df['price'], hist=False, color="r", label="Actual Value")
```


```

sns.distplot(Yhat, hist=False, color="b", label="Fitted Values" , ax=ax1)

plt.title('Actual vs Fitted Values for Price')
plt.xlabel('Price (in dollars)')
plt.ylabel('Proportion of Cars')

plt.show()
plt.close()

```



We can see that the fitted values are reasonably close to the actual values, since the two distributions overlap a bit. However, there is definitely some room for improvement.

### Part 3: Polynomial Regression and Pipelines

Polynomial regression is a particular case of the general linear regression model or multiple linear regression models.

We get non-linear relationships by squaring or setting higher-order terms of the predictor variables.

There are different orders of polynomial regression:

Quadratic - 2nd order

$$\hat{Y} = a + b_1X + b_2X^2$$

Cubic - 3rd order

$$\hat{Y} = a + b_1X + b_2X^2 + b_3X^3$$

Higher order:

$$Y = a + b_1X + b_2X^2 + b_3X^3 \dots$$

We saw earlier that a linear model did not provide the best fit while using highway-mpg as the predictor variable. Let's see if we can try fitting a polynomial model to the data instead.

We will use the following function to plot the data:

```
[44]: def PlotPolly(model, independent_variable, dependent_variabble, Name):
    x_new = np.linspace(15, 55, 100)
    y_new = model(x_new)

    plt.plot(independent_variable, dependent_variabble, '.', x_new, y_new, '-')
    plt.title('Polynomial Fit with Matplotlib for Price ~ Length')
    ax = plt.gca()
    ax.set_facecolor((0.898, 0.898, 0.898))
    fig = plt.gcf()
    plt.xlabel(Name)
    plt.ylabel('Price of Cars')

    plt.show()
    plt.close()
```

lets get the variables

```
[45]: x = df['highway-mpg']
      y = df['price']
```


Let's fit the polynomial using the function polyfit, then use the function poly1d to display the polynomial function.

```
[46]: # Here we use a polynomial of the 3rd order (cubic)
      f = np.polyfit(x, y, 3)
      p = np.poly1d(f)
      print(p)
```

```
-1.557 x3 + 204.8 x2 - 8965 x + 1.379e+05
```

Let's plot the function

```
[47]: PlotPolly(p, x, y, 'highway-mpg')
```



```
[48]: np.polyfit(x, y, 3)
```

```
[48]: array([-1.55663829e+00,  2.04754306e+02, -8.96543312e+03,  1.37923594e+05])
```


We can already see from plotting that this polynomial model performs better than the linear model. This is because the generated polynomial function “hits” more of the data points.

Question #4:

Create 11 order polynomial model with the variables x and y from above?

```
[49]: f1 = np.polyfit(x, y, 11)
      p1 = np.poly1d(f1)
      print(p)
      PlotPolly(p1,x,y, 'Highway MPG')
```

```
          3          2
-1.557 x + 204.8 x - 8965 x + 1.379e+05
```



Double-click here for the solution.

The analytical expression for Multivariate Polynomial function gets complicated. For example, the expression for a second-order (degree=2) polynomial with two variables is given by:

$$\hat{Y} = a + b_1X_1 + b_2X_2 + b_3X_1X_2 + b_4X_1^2 + b_5X_2^2$$

We can perform a polynomial transform on multiple features. First, we import the module:

```
[50]: from sklearn.preprocessing import PolynomialFeatures
```

We create a PolynomialFeatures object of degree 2:

```
[51]: pr=PolynomialFeatures(degree=2)
pr
```

```
[51]: PolynomialFeatures(degree=2, include_bias=True, interaction_only=False)
```

```
[52]: Z_pr=pr.fit_transform(Z)
```

The original data is of 201 samples and 4 features

```
[53]: Z.shape
```



```
[53]: (201, 4)
```

after the transformation, there 201 samples and 15 features

```
[54]: Z_pr.shape
```

```
[54]: (201, 15)
```

Pipeline

Data Pipelines simplify the steps of processing the data. We use the module Pipeline to create a pipeline. We also use StandardScaler as a step in our pipeline.

```
[55]: from sklearn.pipeline import Pipeline
      from sklearn.preprocessing import StandardScaler
```

We create the pipeline, by creating a list of tuples including the name of the model or estimator and its corresponding constructor.

```
[56]: Input=[('scale',StandardScaler()), ('polynomial',PolynomialFeatures(include_bias=False)), ('model',LinearRegression())]
```

we input the list as an argument to the pipeline constructor

```
[57]: pipe=Pipeline(Input)
      pipe
```

```
[57]: Pipeline(memory=None,
      steps=[('scale', StandardScaler(copy=True, with_mean=True, with_std=True)),
      ('polynomial', PolynomialFeatures(degree=2, include_bias=False,
      interaction_only=False)), ('model', LinearRegression(copy_X=True,
      fit_intercept=True, n_jobs=None,
      normalize=False))])
```

We can normalize the data, perform a transform and fit the model simultaneously.

```
[58]: pipe.fit(Z,y)
```

```
/home/jupyterlab/conda/envs/python/lib/python3.6/site-
packages/sklearn/preprocessing/data.py:625: DataConversionWarning: Data with
input dtype int64, float64 were all converted to float64 by StandardScaler.
    return self.partial_fit(X, y)
/home/jupyterlab/conda/envs/python/lib/python3.6/site-
packages/sklearn/base.py:465: DataConversionWarning: Data with input dtype
int64, float64 were all converted to float64 by StandardScaler.
    return self.fit(X, y, **fit_params).transform(X)
```

```
[58]: Pipeline(memory=None,
      steps=[('scale', StandardScaler(copy=True, with_mean=True, with_std=True)),
```

```
(('polynomial', PolynomialFeatures(degree=2, include_bias=False,
interaction_only=False)), ('model', LinearRegression(copy_X=True,
fit_intercept=True, n_jobs=None,
normalize=False)))]
```

Similarly, we can normalize the data, perform a transform and produce a prediction simultaneously

```
[59]: ypipe=pipe.predict(Z)
ypipe[0:4]
```

```
/home/jupyterlab/conda/envs/python/lib/python3.6/site-
packages/sklearn/pipeline.py:331: DataConversionWarning: Data with input dtype
int64, float64 were all converted to float64 by StandardScaler.
Xt = transform.transform(Xt)
```

```
[59]: array([13102.74784201, 13102.74784201, 18225.54572197, 10390.29636555])
```

Question #5:

Create a pipeline that Standardizes the data, then perform prediction using a linear regression model using the features Z and targets y

```
[60]: Input=[('scale',StandardScaler()),('model',LinearRegression())]

pipe=Pipeline(Input)

pipe.fit(Z,y)

ypipe=pipe.predict(Z)
ypipe[0:10]
```

```
/home/jupyterlab/conda/envs/python/lib/python3.6/site-
packages/sklearn/preprocessing/data.py:625: DataConversionWarning: Data with
input dtype int64, float64 were all converted to float64 by StandardScaler.
return self.partial_fit(X, y)
/home/jupyterlab/conda/envs/python/lib/python3.6/site-
packages/sklearn/base.py:465: DataConversionWarning: Data with input dtype
int64, float64 were all converted to float64 by StandardScaler.
return self.fit(X, y, **fit_params).transform(X)
/home/jupyterlab/conda/envs/python/lib/python3.6/site-
packages/sklearn/pipeline.py:331: DataConversionWarning: Data with input dtype
int64, float64 were all converted to float64 by StandardScaler.
Xt = transform.transform(Xt)
```

```
[60]: array([13699.11161184, 13699.11161184, 19051.65470233, 10620.36193015,
15521.31420211, 13869.66673213, 15456.16196732, 15974.00907672,
17612.35917161, 10722.32509097])
```

Double-click here for the solution.

#### Part 4: Measures for In-Sample Evaluation

When evaluating our models, not only do we want to visualize the results, but we also want a quantitative measure to determine how accurate the model is.

Two very important measures that are often used in Statistics to determine the accuracy of a model are:

$R^2$  / R-squared

Mean Squared Error (MSE)

R-squared

R squared, also known as the coefficient of determination, is a measure to indicate how close the data is to the fitted regression line.

The value of the R-squared is the percentage of variation of the response variable (y) that is explained by a linear model.

Mean Squared Error (MSE)

The Mean Squared Error measures the average of the squares of errors, that is, the difference between actual value (y) and the estimated value ( $\hat{y}$ ).

Model 1: Simple Linear Regression

Let's calculate the  $R^2$

```
[61]: #highway_mpg_fit
      lm.fit(X, Y)
      # Find the  $R^2$ 
      print('The R-square is: ', lm.score(X, Y))
```

The R-square is: 0.4965911884339176

We can say that ~ 49.659% of the variation of the price is explained by this simple linear model “horsepower\_fit”.

Let's calculate the MSE

We can predict the output i.e., “yhat” using the predict method, where X is the input variable:

```
[62]: Yhat=lm.predict(X)
      print('The output of the first four predicted value is: ', Yhat[0:4])
```

The output of the first four predicted value is: [16236.50464347 16236.50464347 17058.23802179 13771.3045085 ]

lets import the function mean\_squared\_error from the module metrics

```
[63]: from sklearn.metrics import mean_squared_error
```

we compare the predicted results with the actual results

```
[64]: mse = mean_squared_error(df['price'], Yhat)
print('The mean square error of price and predicted value is: ', mse)
```

The mean square error of price and predicted value is: 31635042.944639888

Model 2: Multiple Linear Regression

Let's calculate the  $R^2$

```
[65]: # fit the model
lm.fit(Z, df['price'])
# Find the  $R^2$ 
print('The R-square is: ', lm.score(Z, df['price']))
```

The R-square is: 0.8093562806577457

We can say that ~ 80.896 % of the variation of price is explained by this multiple linear regression "multi\_fit".

Let's calculate the MSE

we produce a prediction

```
[66]: Y_predict_multifit = lm.predict(Z)
```

we compare the predicted results with the actual results

```
[67]: print('The mean square error of price and predicted value using multifit is: ',
    ↪\
    mean_squared_error(df['price'], Y_predict_multifit))
```

The mean square error of price and predicted value using multifit is:  
11980366.87072649

Model 3: Polynomial Fit

Let's calculate the  $R^2$

let's import the function `r2_score` from the module `metrics` as we are using a different function

```
[68]: from sklearn.metrics import r2_score
```

We apply the function to get the value of  $r^2$

```
[69]: r_squared = r2_score(y, p(x))
print('The R-square value is: ', r_squared)
```

The R-square value is: 0.674194666390652

We can say that ~ 67.419 % of the variation of price is explained by this polynomial fit

MSE

We can also calculate the MSE:

```
[70]: mean_squared_error(df['price'], p(x))
```

```
[70]: 20474146.426361218
```

## Part 5: Prediction and Decision Making

### Prediction

In the previous section, we trained the model using the method fit. Now we will use the method predict to produce a prediction. Lets import pyplot for plotting; we will also be using some functions from numpy.

```
[71]: import matplotlib.pyplot as plt
import numpy as np

%matplotlib inline
```

Create a new input

```
[72]: new_input=np.arange(1, 100, 1).reshape(-1, 1)
```

Fit the model

```
[73]: lm.fit(X, Y)
lm
```

```
[73]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
normalize=False)
```


Produce a prediction

```
[74]: yhat=lm.predict(new_input)
yhat[0:5]
```

```
[74]: array([37601.57247984, 36779.83910151, 35958.10572319, 35136.37234487,
34314.63896655])
```

we can plot the data

```
[75]: plt.plot(new_input, yhat)
plt.show()
```



#### Decision Making: Determining a Good Model Fit

Now that we have visualized the different models, and generated the R-squared and MSE values for the fits, how do we determine a good model fit?

What is a good R-squared value?

When comparing models, the model with the higher R-squared value is a better fit for the data.

What is a good MSE?

When comparing models, the model with the smallest MSE value is a better fit for the data.

Let's take a look at the values for the different models.

Simple Linear Regression: Using Highway-mpg as a Predictor Variable of Price.

R-squared: 0.49659118843391759

MSE:  $3.16 \times 10^7$

Multiple Linear Regression: Using Horsepower, Curb-weight, Engine-size, and Highway-mpg as Predictor Variables of Price.

R-squared: 0.80896354913783497

MSE:  $1.2 \times 10^7$

Polynomial Fit: Using Highway-mpg as a Predictor Variable of Price.

R-squared: 0.6741946663906514

MSE:  $2.05 \times 10^7$

## Simple Linear Regression model (SLR) vs Multiple Linear Regression model (MLR)

Usually, the more variables you have, the better your model is at predicting, but this is not always true. Sometimes you may not have enough data, you may run into numerical problems, or many of the variables may not be useful and or even act as noise. As a result, you should always check the MSE and  $R^2$ .

So to be able to compare the results of the MLR vs SLR models, we look at a combination of both the R-squared and MSE to make the best conclusion about the fit of the model.

MSE The MSE of SLR is  $3.16 \times 10^7$  while MLR has an MSE of  $1.2 \times 10^7$ . The MSE of MLR is much smaller.

R-squared: In this case, we can also see that there is a big difference between the R-squared of the SLR and the R-squared of the MLR. The R-squared for the SLR ( $\sim 0.497$ ) is very small compared to the R-squared for the MLR ( $\sim 0.809$ ).

This R-squared in combination with the MSE show that MLR seems like the better model fit in this case, compared to SLR.

## Simple Linear Model (SLR) vs Polynomial Fit

MSE: We can see that Polynomial Fit brought down the MSE, since this MSE is smaller than the one from the SLR.

R-squared: The R-squared for the Polyfit is larger than the R-squared for the SLR, so the Polynomial Fit also brought up the R-squared quite a bit.

Since the Polynomial Fit resulted in a lower MSE and a higher R-squared, we can conclude that this was a better fit model than the simple linear regression for predicting Price with Highway-mpg as a predictor variable.

## Multiple Linear Regression (MLR) vs Polynomial Fit

MSE: The MSE for the MLR is smaller than the MSE for the Polynomial Fit.

R-squared: The R-squared for the MLR is also much larger than for the Polynomial Fit.

## Conclusion:

Comparing these three models, we conclude that the MLR model is the best model to be able to predict price from our dataset. This result makes sense, since we have 27 variables in total, and we know that more than one of those variables are potential predictors of the final car price.

Thank you for completing this notebook

[<p><a href="https://cocl.us/corsera\\_da0101en\\_notebook\\_bottom"><img src="https://s3-api.us-geo.](https://cocl.us/corsera_da0101en_notebook_bottom)

## About the Authors:

This notebook was written by Mahdi Noorian PhD, Joseph Santarcangelo, Bahare Talayian, Eric Xiao, Steven Dong, Parizad, Hima Vsudevan and Fiorella Wenver and Yi Yao.

Joseph Santarcangelo is a Data Scientist at IBM, and holds a PhD in Electrical Engineering. His research focused on using Machine Learning, Signal Processing, and Computer Vision to determine how videos impact human cognition. Joseph has been working for IBM since he completed his PhD.

Copyright © 2018 IBM Developer Skills Network. This notebook and its source code are released under the terms of the MIT License.